

PRACTICAL 1

➤ Source Code

```
import pandas as pd
import numpy as np

df = pd.read_csv("prac1_sample.csv")

DIVIDER = "=" * 60
print(f"\n{DIVIDER}")
print(" STEP 0 — Raw Data")
print(f" Shape: {df.shape}")
print(DIVIDER)
print(df.to_string(index=False))

print(f"\n{DIVIDER}")
print(" STEP 1 — Missing Value Count")
print(DIVIDER)
print(df.isnull().sum())

print(f"\n{DIVIDER}")
print(" STEP 2 — Statistical Summary (describe)")
print(DIVIDER)
print(df.describe())

print(f"\n{DIVIDER}")
print(" STEP 3 — Data Types")
print(DIVIDER)
print(df.dtypes)

df = df.drop(['Cabin', 'PassengerId', 'Name', 'Ticket'], axis=1)

print(f"\n{DIVIDER}")
print(" STEP 4 — After Dropping Cabin, PassengerId, Name, Ticket")
print(f" Shape: {df.shape}")
print(DIVIDER)
print(df.to_string(index=False))

df['Age'] = df['Age'].fillna(df['Age'].mean())
df['Embarked'] = df['Embarked'].fillna(df['Embarked'].mode()[0])

print(f"\n{DIVIDER}")
```

```

print(" STEP 5 — After Filling Missing Values")
print(" Age filled with mean | Embarked filled with mode")
print(DIVIDER)
print(df.to_string(index=False))

df['Fare'] = (df['Fare'] - df['Fare'].min()) / (df['Fare'].max() - df['Fare'].min())

print(f"\n{DIVIDER}")
print(" STEP 6 — After Normalizing Fare column (min-max)")
print(DIVIDER)
print(df.to_string(index=False))

df['Age'] = df['Age'].astype(float)

print(f"\n{DIVIDER}")
print(" STEP 7 — Age converted to float")
print(DIVIDER)
print(df.dtypes)

df['Sex'] = df['Sex'].map({'male': 0, 'female': 1})

print(f"\n{DIVIDER}")
print(" STEP 8 — After Label Encoding: Sex (male=0, female=1)")
print(DIVIDER)
print(df.to_string(index=False))

df = pd.get_dummies(df, columns=['Embarked'])

print(f"\n{DIVIDER}")
print(" STEP 9 — FINAL: After One-Hot Encoding Embarked column")
print(f" Shape: {df.shape}")
print(DIVIDER)
print(df.to_string(index=False))

print(f"\n{'='*60}")
print(" PRACTICAL 1 COMPLETE")
print(f"\n{'='*60}")

```

➤ Output

```
Ubuntu
(venv) shreehari@DESKTOP-0F0T52N: /mnt/d/Third Year/DSBDA/practical1$ python3 prac_1.py

=====
STEP 0 - Raw Data
Shape: (6, 12)
=====
PassengerId  Survived  Pclass      Name      Sex  Age  SibSp  Parch  Ticket  Fare  Cabin  Embarked
1            0         3  Braund, Mr. Owen  male  22.0    1     0  A/5 21171  7.25   NaN     S
2            1         1  Cumings, Mrs. John female  38.0    1     0  PC 17599  71.28  C85     C
3            1         3  Heikkinen, Miss. Laina female  26.0    0     0  STON/O2  7.92   NaN     S
4            1         1  Futrelle, Mrs. Jacques female  35.0    1     0  113803  53.10  C123    S
5            0         3  Allen, Mr. William male   NaN    0     0  373450  8.05   NaN     NaN
6            0         3  Moran, Mr. James  male   NaN    0     0  330877  8.46   NaN     Q

=====
STEP 1 - Missing Value Count
=====
PassengerId  0
Survived      0
Pclass        0
Name          0
Sex           0
Age           2
SibSp         0
Parch         0
Ticket        0
Fare          0
Cabin         4
Embarked      1
dtype: int64

=====
```

```
Ubuntu
STEP 2 - Statistical Summary (describe)
=====
count      PassengerId  Survived  Pclass    Age    SibSp  Parch    Fare
mean         3.500000    0.500000    2.333333  30.25  0.500000  0.0  26.010000
std          1.870829    0.547723    1.032796   7.50  0.547723  0.0  28.611154
min          1.000000    0.000000    1.000000  22.00  0.000000  0.0   7.250000
25%          2.250000    0.000000    1.500000  25.00  0.000000  0.0   7.952500
50%          3.500000    0.500000    3.000000  30.50  0.500000  0.0   8.255000
75%          4.750000    1.000000    3.000000  35.75  1.000000  0.0  41.940000
max          6.000000    1.000000    3.000000  38.00  1.000000  0.0  71.280000

=====
STEP 3 - Data Types
=====
PassengerId    int64
Survived        int64
Pclass          int64
Name            str
Sex             str
Age            float64
SibSp           int64
Parch           int64
Ticket          str
Fare            float64
Cabin           str
Embarked        str
dtype: object

=====
```

```
Ubuntu
STEP 4 - After Dropping Cabin, PassengerId, Name, Ticket
Shape: (6, 8)
=====
Survived  Pclass  Sex   Age   SibSp  Parch  Fare  Embarked
0         3     male  22.0   1      0    7.25    S
1         1     female 38.0   1      0   71.28    C
1         3     female 26.0   0      0    7.92    S
1         1     female 35.0   1      0   53.10    S
0         3     male   NaN    0      0    8.05    NaN
0         3     male   NaN    0      0    8.46    Q
=====

STEP 5 - After Filling Missing Values
Age filled with mean | Embarked filled with mode
=====
Survived  Pclass  Sex   Age   SibSp  Parch  Fare  Embarked
0         3     male  22.00   1      0    7.25    S
1         1     female 38.00   1      0   71.28    C
1         3     female 26.00   0      0    7.92    S
1         1     female 35.00   1      0   53.10    S
0         3     male  30.25   0      0    8.05    S
0         3     male  30.25   0      0    8.46    Q
=====

STEP 6 - After Normalizing Fare column (min-max)
=====
Survived  Pclass  Sex   Age   SibSp  Parch  Fare  Embarked
0         3     male  22.00   1      0  0.000000    S
1         1     female 38.00   1      0  1.000000    C
1         3     female 26.00   0      0  0.010464    S
1         1     female 35.00   1      0  0.716071    S
0         3     male  30.25   0      0  0.012494    S
=====

Ubuntu
6:36 PM
4/23/2026
```

```
Ubuntu
STEP 7 - Age converted to float
=====
Survived    int64
Pclass      int64
Sex         str
Age         float64
SibSp       int64
Parch       int64
Fare        float64
Embarked    str
dtype: object
=====

STEP 8 - After Label Encoding: Sex (male=0, female=1)
=====
Survived  Pclass  Sex   Age   SibSp  Parch  Fare  Embarked
0         3     0  22.00   1      0  0.000000    S
1         1     1  38.00   1      0  1.000000    C
1         3     1  26.00   0      0  0.010464    S
1         1     1  35.00   1      0  0.716071    S
0         3     0  30.25   0      0  0.012494    S
0         3     0  30.25   0      0  0.01897    Q
=====

STEP 9 - FINAL: After One-Hot Encoding Embarked column
Shape: (6, 10)
=====
Survived  Pclass  Sex   Age   SibSp  Parch  Fare  Embarked_C  Embarked_Q  Embarked_S
0         3     0  22.00   1      0  0.000000    False    False    True
1         1     1  38.00   1      0  1.000000    True     False    False
=====

Ubuntu
6:36 PM
4/23/2026
```

```
=====
STEP 8 - After Label Encoding: Sex (male=0, female=1)
=====
Survived  Pclass  Sex    Age    SibSp  Parch    Fare  Embarked
0         3      0    22.00    1      0  0.000000    S
1         1      1    38.00    1      0  1.000000    C
1         3      1    26.00    0      0  0.010464    S
1         1      1    35.00    1      0  0.716071    S
0         3      0    30.25    0      0  0.012494    S
0         3      0    30.25    0      0  0.018897    Q

=====
STEP 9 - FINAL: After One-Hot Encoding Embarked column
Shape: (6, 10)
=====
Survived  Pclass  Sex    Age    SibSp  Parch    Fare  Embarked_C  Embarked_Q  Embarked_S
0         3      0    22.00    1      0  0.000000    False      False      True
1         1      1    38.00    1      0  1.000000    True       False      False
1         3      1    26.00    0      0  0.010464    False     False      True
1         1      1    35.00    1      0  0.716071    False     False      True
0         3      0    30.25    0      0  0.012494    False     False      True
0         3      0    30.25    0      0  0.018897    False     True       False

=====
PRACTICAL 1 COMPLETE
=====

(venv) shreehari@DESKTOP-0FDT52N: /mnt/d/Third Year/D58DA/practical1$
```

PRACTICAL 2

➤ Source Code

```
import pandas as pd
import numpy as np
DIVIDER = "=" * 60
df = pd.read_csv("prac2_samples.csv")
print(f"\n{DIVIDER}")
print(" STEP 0 — Raw Dataset (contains bad values)")
print(f" Shape: {df.shape}")
print(DIVIDER)
print(df.to_string(index=False))
print("\n Notice: Maths has 150 and -5 (impossible marks)")
print("      Science has 300 (impossible)")
print("      Some cells are NaN (missing)")

print(f"\n{DIVIDER}")
print(" STEP 1 — Missing Value Count")
print(DIVIDER)
print(df.isnull().sum())

print(f"\n{DIVIDER}")
print(" STEP 2 — Statistical Summary")
print(DIVIDER)
print(df.describe())
df['Maths'] = df['Maths'].apply(lambda x: np.nan if pd.isna(x) and (x < 0 or x > 100) else x)
df['Science'] = df['Science'].apply(lambda x: np.nan if pd.isna(x) and x > 100 else x)

print(f"\n{DIVIDER}")
print(" STEP 3 — After Fixing Inconsistencies")
print(" (marks < 0 or > 100 replaced with NaN)")
print(DIVIDER)
print(df.to_string(index=False))
print(f"\n Missing values now: {df.isnull().sum()}")

df.fillna(df.mean(numeric_only=True), inplace=True)

print(f"\n{DIVIDER}")
print(" STEP 4 — After Filling NaN with Column Mean")
print(DIVIDER)
print(df.round(2).to_string(index=False))
print(f"\n Missing values now: {df.isnull().sum()}")
```

```

cols = ['Maths', 'Science', 'English']
Q1 = df[cols].quantile(0.25)
Q3 = df[cols].quantile(0.75)
IQR = Q3 - Q1

print(f"\n{DIVIDER}")
print(" STEP 5 — Outlier Detection (IQR Method)")
print(DIVIDER)
print(f" Q1:\n{Q1.round(2)}\n")
print(f" Q3:\n{Q3.round(2)}\n")
print(f" IQR:\n{IQR.round(2)}\n")

lower = Q1 - 1.5 * IQR
upper = Q3 + 1.5 * IQR
print(f" Lower bound:\n{lower.round(2)}\n")
print(f" Upper bound:\n{upper.round(2)}\n")
outlier_mask = ((df[cols] < lower) | (df[cols] > upper)).any(axis=1)
print(f"\n Rows identified as outliers:\n{df[outlier_mask].round(2)}")
df_clean = df[~outlier_mask].copy()
print(f"\n Shape before removing outliers: {df.shape}")
print(f" Shape after removing outliers: {df_clean.shape}")
df_clean['Maths_norm'] = (
    (df_clean['Maths'] - df_clean['Maths'].min()) /
    (df_clean['Maths'].max() - df_clean['Maths'].min())
)
print(f"\n{DIVIDER}")
print(" STEP 6 — After Normalizing Maths column (0 to 1 scale)")
print(DIVIDER)
print(df_clean[['Student_ID', 'Maths', 'Maths_norm']].round(4).to_string(index=False))
df_clean['Science_log'] = np.log(df_clean['Science'])
print(f"\n{DIVIDER}")
print(" STEP 7 — After Log Transformation on Science column")
print(" (reduces skewness, good for right-skewed data)")
print(DIVIDER)
print(df_clean[['Student_ID', 'Science', 'Science_log']].round(4).to_string(index=False))
print(f"\n{DIVIDER}")
print(" FINAL — Complete Cleaned & Transformed Dataset")
print(f" Shape: {df_clean.shape}")
print(DIVIDER)
print(df_clean.round(4).to_string(index=False))

print(f"\n{'='*60}")
print(" PRACTICAL 2 COMPLETE")
print(f"\n{'='*60}")

```

➤ Output

```
Ubuntu
x + v
(venv) shreehari@DESKTOP-0F0T52N:/mnt/d/Third Year/DSBDA/practical:$ python3 prac_2.py

=====
STEP 0 - Raw Dataset (contains bad values)
Shape: (10, 4)
=====
Student_ID  Maths  Science  English
1    85.0    80.0    78.0
2    92.0    85.0    82.0
3     NaN    70.0    88.0
4    45.0     NaN    90.0
5    78.0    75.0     NaN
6   150.0    95.0    85.0
7    88.0    89.0    87.0
8    76.0   300.0    92.0
9    -5.0    65.0    79.0
10   95.0    88.0    84.0

Notice: Maths has 150 and -5 (impossible marks)
Science has 300 (impossible)
Some cells are NaN (missing)

=====
STEP 1 - Missing Value Count
=====
Student_ID    0
Maths         1
Science       1
English       1
dtype: int64
=====

=====
STEP 2 - Statistical Summary
=====
Student_ID  Maths  Science  English
count    10.00000  9.000000  9.000000  9.000000
mean      5.50000  78.222222  105.222222  85.000000
std       3.02765  41.532450  73.671191  4.769696
min       1.00000  -5.000000  65.000000  78.000000
25%       3.25000  76.000000  75.000000  82.000000
50%       5.50000  85.000000  85.000000  85.000000
75%       7.75000  92.000000  89.000000  88.000000
max      10.00000  150.000000  300.000000  92.000000

=====
STEP 3 - After Fixing Inconsistencies
(marks < 0 or > 100 replaced with NaN)
=====
Student_ID  Maths  Science  English
1    85.0    80.0    78.0
2    92.0    85.0    82.0
3     NaN    70.0    88.0
4    45.0     NaN    90.0
5    78.0    75.0     NaN
6     NaN    95.0    85.0
7    88.0    89.0    87.0
8    76.0     NaN    92.0
9     NaN    65.0    79.0
10   95.0    88.0    84.0

Missing values now:
Student_ID    0
```

```
Ubuntu
x + v
(venv) shreehari@DESKTOP-0F0T52N:/mnt/d/Third Year/DSBDA/practical:$ python3 prac_2.py

=====
STEP 0 - Raw Dataset (contains bad values)
Shape: (10, 4)
=====
Student_ID  Maths  Science  English
1    85.0    80.0    78.0
2    92.0    85.0    82.0
3     NaN    70.0    88.0
4    45.0     NaN    90.0
5    78.0    75.0     NaN
6   150.0    95.0    85.0
7    88.0    89.0    87.0
8    76.0   300.0    92.0
9    -5.0    65.0    79.0
10   95.0    88.0    84.0

Notice: Maths has 150 and -5 (impossible marks)
Science has 300 (impossible)
Some cells are NaN (missing)

=====
STEP 1 - Missing Value Count
=====
Student_ID    0
Maths         1
Science       1
English       1
dtype: int64
=====

=====
STEP 2 - Statistical Summary
=====
Student_ID  Maths  Science  English
count    10.00000  9.000000  9.000000  9.000000
mean      5.50000  78.222222  105.222222  85.000000
std       3.02765  41.532450  73.671191  4.769696
min       1.00000  -5.000000  65.000000  78.000000
25%       3.25000  76.000000  75.000000  82.000000
50%       5.50000  85.000000  85.000000  85.000000
75%       7.75000  92.000000  89.000000  88.000000
max      10.00000  150.000000  300.000000  92.000000

=====
STEP 3 - After Fixing Inconsistencies
(marks < 0 or > 100 replaced with NaN)
=====
Student_ID  Maths  Science  English
1    85.0    80.0    78.0
2    92.0    85.0    82.0
3     NaN    70.0    88.0
4    45.0     NaN    90.0
5    78.0    75.0     NaN
6     NaN    95.0    85.0
7    88.0    89.0    87.0
8    76.0     NaN    92.0
9     NaN    65.0    79.0
10   95.0    88.0    84.0

Missing values now:
Student_ID    0
```



```
Ubuntu x + v

Missing values now:
Student_ID 0
Maths      3
Science     2
English     1
dtype: int64

=====
STEP 4 - After Filling NaN with Column Mean
=====
Student_ID Maths Science English
1 85.00 80.00 78.0
2 92.00 85.00 82.0
3 79.86 70.00 88.0
4 45.00 80.88 90.0
5 78.00 75.00 85.0
6 79.86 95.00 85.0
7 88.00 89.00 87.0
8 76.00 80.88 92.0
9 79.86 65.00 79.0
10 95.00 88.00 84.0

Missing values now:
Student_ID 0
Maths      0
Science     0
English     0
dtype: int64

=====
STEP 5 - Outlier Detection (IQR Method)
```

```
Ubuntu x + v

=====
STEP 5 - Outlier Detection (IQR Method)
=====
Q1:
Maths      78.46
Science    76.25
English    82.50
Name: 0.25, dtype: float64

Q3:
Maths      87.25
Science    87.25
English    87.75
Name: 0.75, dtype: float64

IQR:
Maths      8.79
Science    11.00
English     5.25
dtype: float64

Lower bound:
Maths      65.29
Science    59.75
English    74.62
dtype: float64

Upper bound:
Maths     100.43
Science   103.75
English    95.62
dtype: float64
```

```
Ubuntu x + v
Rows identified as outliers:
Student_ID Maths Science English
3 4 45.0 80.88 90.0

Shape before removing outliers: (10, 4)
Shape after removing outliers: (9, 4)

=====
STEP 6 - After Normalizing Maths column (0 to 1 scale)
=====
Student_ID Maths Maths_norm
1 85.0000 0.4737
2 92.0000 0.8421
3 79.8571 0.2030
5 78.0000 0.1053
6 79.8571 0.2030
7 88.0000 0.6316
8 76.0000 0.0000
9 79.8571 0.2030
10 95.0000 1.0000

=====
STEP 7 - After Log Transformation on Science column
(reduces skewness, good for right-skewed data)
=====
Student_ID Science Science_log
1 80.000 4.3820
2 85.000 4.4427
3 70.000 4.2485
5 75.000 4.3175
6 95.000 4.5539
```

```
Ubuntu x + v
=====
Student_ID Science Science_log
1 80.000 4.3820
2 85.000 4.4427
3 70.000 4.2485
5 75.000 4.3175
6 95.000 4.5539
7 89.000 4.4886
8 80.875 4.3929
9 65.000 4.1744
10 88.000 4.4773

=====
FINAL - Complete Cleaned & Transformed Dataset
Shape: (9, 6)
=====
Student_ID Maths Science English Maths_norm Science_log
1 85.0000 80.000 78.0 0.4737 4.3820
2 92.0000 85.000 82.0 0.8421 4.4427
3 79.8571 70.000 88.0 0.2030 4.2485
5 78.0000 75.000 85.0 0.1053 4.3175
6 79.8571 95.000 85.0 0.2030 4.5539
7 88.0000 89.000 87.0 0.6316 4.4886
8 76.0000 80.875 92.0 0.0000 4.3929
9 79.8571 65.000 79.0 0.2030 4.1744
10 95.0000 88.000 84.0 1.0000 4.4773

=====
PRACTICAL 2 COMPLETE
=====
(venv) shreehari@DESKTOP-0FOT52N: /mnt/d/Third Year/DS80A/practicals$
```

PRACTICAL 3

➤ Source Code

```
import pandas as pd
import numpy as np
DIVIDER = "=" * 60
print(f"\n{'#' * 60}")
print(" PART A — Income Dataset (Grouped Statistics)")
print(f"{'#' * 60}")
df = pd.read_csv("income_data.csv")

print(f"\n{DIVIDER}")
print(" STEP 0 — Raw Income Dataset")
print(DIVIDER)
print(df.to_string(index=False))

print(f"\n{DIVIDER}")
print(" STEP 1 — Summary Statistics Grouped by Age_Group")
print(DIVIDER)
summary = df.groupby('Age_Group')['Income'].agg(['mean', 'median', 'min', 'max', 'std'])
summary.columns = ['Mean', 'Median', 'Min', 'Max', 'Std Dev']
print(summary.round(2))

print(f"\n{DIVIDER}")
print(" STEP 2 — Income Values Listed per Group")
print(DIVIDER)
grouped_list = df.groupby('Age_Group')['Income'].apply(list)
for group, values in grouped_list.items():
    print(f" {group}: {values}")
print(f"\n{DIVIDER}")
print(" STEP 3 — Overall Income Statistics")
print(DIVIDER)
print(f" Mean : {df['Income'].mean():.2f}")
print(f" Median : {df['Income'].median():.2f}")
print(f" Mode : {df['Income'].mode().values}")
print(f" Std Dev: {df['Income'].std():.2f}")
print(f" Min : {df['Income'].min():.2f}")
print(f" Max : {df['Income'].max():.2f}")

print(f"\n\n{'#' * 60}")
```

```

print(" PART B — Iris Dataset (Species-wise Statistics)")
print(f"{'#'*60}")
iris = pd.read_csv("iris_data.csv")

print(f"\n{DIVIDER}")
print(" STEP 0 — Iris Dataset")
print(DIVIDER)
print(iris.to_string(index=False))

setosa = iris[iris['species'] == 'Iris-setosa']
versicolor = iris[iris['species'] == 'Iris-versicolor']
virginica = iris[iris['species'] == 'Iris-virginica']
print(f"\n{DIVIDER}")
print(" STEP 1 — Setosa: describe()")
print(DIVIDER)
print(setosa.describe().round(3))
print(f"\n{DIVIDER}")
print(" STEP 2 — Versicolor: describe()")
print(DIVIDER)
print(versicolor.describe().round(3))

print(f"\n{DIVIDER}")
print(" STEP 3 — Virginica: describe()")
print(DIVIDER)
print(virginica.describe().round(3))

print(f"\n{DIVIDER}")
print(" STEP 4 — Percentiles for Sepal Length (all species)")
print(DIVIDER)
for name, grp in [('Setosa', setosa), ('Versicolor', versicolor), ('Virginica', virginica)]:
    p = np.percentile(grp['sepal_length'], [25, 50, 75])
    print(f" {name:12s} → 25th: {p[0]:.2f} | 50th (Median): {p[1]:.2f} | 75th: {p[2]:.2f}")
print(f"\n{DIVIDER}")
print(" STEP 5 — Mean of All Features by Species")
print(DIVIDER)
print(iris.groupby('species').mean().round(3))

print(f"\n{'='*60}")
print(" PRACTICAL 3 COMPLETE")
print(f"{'='*60}\n")

```

➤ Output

```
Ubuntu
(venv) shreehari@DESKTOP-0FOT52N: /mnt/d/Third Year/DSBDA/practical$ python3 prac_3.py
PART A - Income Dataset (Grouped Statistics)

=====
STEP 0 - Raw Income Dataset
=====
Name Age Income Age_Group
Alice 25 30000 Young
Bob 32 50000 Young
Carol 45 75000 Middle
Dave 23 28000 Young
Eve 55 90000 Senior
Frank 40 65000 Middle
Grace 29 40000 Young
Hank 60 85000 Senior
Iris 35 55000 Middle
Jack 48 70000 Middle

=====
STEP 1 - Summary Statistics Grouped by Age_Group
=====
Age_Group Mean Median Min Max Std Dev
Middle 66250.0 67500.0 55000 75000 8539.13
Senior 87500.0 87500.0 85000 90000 3535.53
Young 37000.0 35000.0 28000 50000 10132.46

=====
STEP 2 - Income Values Listed per Group
=====
Middle: [75000, 65000, 55000, 70000]
Senior: [90000, 85000]
```

```
Ubuntu
=====
STEP 3 - Overall Income Statistics
=====
Mean : 58,800.00
Median : 60,000.00
Mode : [28000 30000 40000 50000 55000 65000 70000 75000 85000 90000]
Std Dev: 21,882.51
Min : 28,000.00
Max : 90,000.00
PART B - Iris Dataset (Species-wise Statistics)

=====
STEP 0 - Iris Dataset
=====
sepal_length sepal_width petal_length petal_width species
5.1 3.5 1.4 0.2 Iris-setosa
4.9 3.0 1.4 0.2 Iris-setosa
4.7 3.2 1.3 0.2 Iris-setosa
7.0 3.2 4.7 1.4 Iris-versicolor
6.4 3.2 4.5 1.5 Iris-versicolor
6.9 3.1 4.9 1.5 Iris-versicolor
6.3 3.3 6.0 2.5 Iris-virginica
5.8 2.7 5.1 1.9 Iris-virginica
7.1 3.0 5.9 2.1 Iris-virginica
6.5 3.2 5.8 2.2 Iris-virginica

=====
STEP 1 - Setosa: describe()
=====
count sepal_length sepal_width petal_length petal_width
count 3.0 3.000 3.000 3.0
```

```
min      4.7      3.000      1.300      0.2
25%      4.8      3.100      1.350      0.2
50%      4.9      3.200      1.400      0.2
75%      5.0      3.350      1.400      0.2
max      5.1      3.500      1.400      0.2

=====
STEP 2 - Versicolor: describe()
=====
      sepal_length  sepal_width  petal_length  petal_width
count      3.000      3.000      3.0      3.000
mean      6.767      3.167      4.7      1.467
std       0.321      0.058      0.2      0.058
min      6.400      3.100      4.5      1.400
25%      6.650      3.150      4.6      1.450
50%      6.900      3.200      4.7      1.500
75%      6.950      3.200      4.8      1.500
max      7.000      3.200      4.9      1.500

=====
STEP 3 - Virginica: describe()
=====
      sepal_length  sepal_width  petal_length  petal_width
count      4.000      4.000      4.000      4.000
mean      6.425      3.050      5.700      2.175
std       0.538      0.265      0.408      0.250
min      5.800      2.700      5.100      1.900
25%      6.175      2.925      5.625      2.050
50%      6.400      3.100      5.850      2.150
75%      6.650      3.225      5.925      2.275
max      7.100      3.300      6.000      2.500

=====
```

```
=====
      sepal_length  sepal_width  petal_length  petal_width
count      4.000      4.000      4.000      4.000
mean      6.425      3.050      5.700      2.175
std       0.538      0.265      0.408      0.250
min      5.800      2.700      5.100      1.900
25%      6.175      2.925      5.625      2.050
50%      6.400      3.100      5.850      2.150
75%      6.650      3.225      5.925      2.275
max      7.100      3.300      6.000      2.500

=====
STEP 4 - Percentiles for Sepal Length (all species)
=====
Setosa      → 25th: 4.80 | 50th (Median): 4.90 | 75th: 5.00
Versicolor → 25th: 6.65 | 50th (Median): 6.90 | 75th: 6.95
Virginica   → 25th: 6.17 | 50th (Median): 6.40 | 75th: 6.65

=====
STEP 5 - Mean of All Features by Species
=====
      sepal_length  sepal_width  petal_length  petal_width
species
Iris-setosa      4.900      3.233      1.367      0.200
Iris-versicolor  6.767      3.167      4.700      1.467
Iris-virginica   6.425      3.050      5.700      2.175

=====
PRACTICAL 3 COMPLETE
=====

(venv) shreehari@DESKTOP-0F0T52N: /mnt/d/Third Year/DSBDA/practicals$
```

PRACTICAL 4

➤ Source Code

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
DIVIDER = "=" * 60
from sklearn.datasets import fetch_california_housing
housing = fetch_california_housing(as_frame=True)
df = housing.frame
df = df.rename(columns={'MedHouseVal': 'MEDV'})
print(f"\n{DIVIDER}")
print(" STEP 1 — Dataset Loaded (California Housing)")
print(f" Shape: {df.shape}")
print(DIVIDER)
print(df.head().to_string(index=False))
print(f"\n{DIVIDER}")
print(" STEP 2 — Dataset Info & Missing Values")
print(DIVIDER)
print(df.info())
print(f"\n Missing values:\n{df.isnull().sum()}")

print(f"\n{DIVIDER}")
print(" STEP 3 — Statistical Summary")
print(DIVIDER)
print(df.describe().round(3))
X = df.drop('MEDV', axis=1) # All columns except target
y = df['MEDV']             # Target: House price
print(f"\n{DIVIDER}")
print(" STEP 4 — Features (X) and Target (y)")
print(DIVIDER)
print(f" Feature columns: {list(X.columns)}")
print(f" Target column : MEDV (Median House Value)")
print(f" X shape: {X.shape} | y shape: {y.shape}")

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
```

```

print(f"\n{DIVIDER}")
print(" STEP 5 — Train-Test Split (80% train, 20% test)")
print(DIVIDER)
print(f" Training set : {X_train.shape[0]} rows")
print(f" Testing set : {X_test.shape[0]} rows")
model = LinearRegression()
model.fit(X_train, y_train)
print(f"\n{DIVIDER}")
print(" STEP 6 — Model Trained!")
print(DIVIDER)
print(f" Intercept (b0): {model.intercept_:.4f}")
print(f"\n Coefficients:")
for feat, coef in zip(X.columns, model.coef_):
    print(f"   {feat:20s}: {coef:.4f}")
y_pred = model.predict(X_test)
print(f"\n{DIVIDER}")
print(" STEP 7 — Sample Predictions vs Actual")
print(DIVIDER)
comparison = pd.DataFrame({
    'Actual Price': y_test.values[:10],
    'Predicted Price': y_pred[:10].round(3)
})
print(comparison.to_string(index=False))

mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)

print(f"\n{DIVIDER}")
print(" STEP 8 — Model Evaluation Metrics")
print(DIVIDER)
print(f" Mean Squared Error (MSE) : {mse:.4f}")
print(f" Root MSE (RMSE): {rmse:.4f}")
print(f" R2 Score : {r2:.4f}")
print(f"\n Interpretation:")
print(f" → R2 = {r2:.4f} means the model explains {r2*100:.1f}% of variance in price.")
print(f" → Closer to 1.0 is better.")

print(f"\n{'='*60}")
print(" PRACTICAL 4 COMPLETE")
print(f"\n{'='*60}")

```


➤ Output

```
Ubuntu
(venv) shreehari@DESKTOP-0F0T52N:/mnt/d/Third Year/DSBDA/practical:$ python3 prac_4.py

=====
STEP 1 - Dataset Loaded (California Housing)
Shape: (20640, 9)
=====
MedInc   HouseAge  AveRooms  AveBedrms  Population  AveOccup  Latitude  Longitude  MEDV
8.3252   41.0      6.984127  1.023810   322.0       2.555556  37.88    -122.23    4.526
8.3014   21.0      6.238137  0.971880   2401.0      2.109842  37.86    -122.22    3.585
7.2574   52.0      8.288136  1.073446   496.0       2.802260  37.85    -122.24    3.521
5.6431   52.0      5.817352  1.073059   558.0       2.547945  37.85    -122.25    3.413
3.8462   52.0      6.281853  1.081081   565.0       2.181467  37.85    -122.25    3.422

=====
STEP 2 - Dataset Info & Missing Values
=====
<class 'pandas.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 9 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   MedInc       20640 non-null  float64
1   HouseAge     20640 non-null  float64
2   AveRooms     20640 non-null  float64
3   AveBedrms    20640 non-null  float64
4   Population   20640 non-null  float64
5   AveOccup     20640 non-null  float64
6   Latitude     20640 non-null  float64
7   Longitude    20640 non-null  float64
8   MEDV         20640 non-null  float64
dtypes: float64(9)
memory usage: 1.4 MB
```

```
Ubuntu
Missing values:
MedInc      0
HouseAge    0
AveRooms    0
AveBedrms   0
Population  0
AveOccup    0
Latitude    0
Longitude   0
MEDV        0
dtype: int64

=====
STEP 3 - Statistical Summary
=====
count    MedInc   HouseAge  AveRooms  AveBedrms  Population  AveOccup  Latitude  Longitude  MEDV
20640.000 20640.000 20640.000 20640.000 20640.000 20640.000 20640.000 20640.000 20640.000 20640.000
mean      3.871    28.639    5.429     1.097     1425.477    3.071     35.632    -119.570    2.069
std       1.900    12.586    2.474     0.474    1132.462    10.386     2.136     2.004     1.154
min       0.500     1.000     0.846     0.333     3.000     0.692     32.540    -124.350    0.150
25%       2.563    18.000    4.441     1.006     787.000    2.430     33.930    -121.800    1.196
50%       3.535    29.000    5.229     1.049    1166.000    2.818     34.260    -118.490    1.797
75%       4.743    37.000    6.052     1.100    1725.000    3.282     37.710    -118.010    2.647
max       15.000    52.000   141.909    34.067   35682.000  1243.333    41.950    -114.310    5.000

=====
STEP 4 - Features (X) and Target (y)
=====
Feature columns: ['MedInc', 'HouseAge', 'AveRooms', 'AveBedrms', 'Population', 'AveOccup', 'Latitude', 'Longitude']
Target column : MEDV (Median House Value)
X shape: (20640, 8) | y shape: (20640,)
```

```
=====  
STEP 5 - Train-Test Split (80% train, 20% test)  
=====  
Training set : 16512 rows  
Testing set  : 4128 rows  
  
=====  
STEP 6 - Model Trained!  
=====  
Intercept (b0): -37.0233  
  
Coefficients:  
MedInc      : 0.4487  
HouseAge    : 0.0097  
AveRooms    : -0.1233  
AveBedrms   : 0.7831  
Population  : -0.0000  
AveOccup    : -0.0035  
Latitude    : -0.4198  
Longitude   : -0.4337  
  
=====  
STEP 7 - Sample Predictions vs Actual  
=====  
Actual Price Predicted Price  
0.47700      0.719  
0.45800      1.764  
5.00001      2.710  
2.18600      2.839  
2.78000      2.605  
1.58700      2.012
```

```
=====  
STEP 7 - Sample Predictions vs Actual  
=====  
Actual Price Predicted Price  
0.47700      0.719  
0.45800      1.764  
5.00001      2.710  
2.18600      2.839  
2.78000      2.605  
1.58700      2.012  
1.98200      2.646  
1.57500      2.169  
3.40000      2.741  
4.46600      3.916  
  
=====  
STEP 8 - Model Evaluation Metrics  
=====  
Mean Squared Error (MSE) : 0.5559  
Root MSE (RMSE) : 0.7456  
R2 Score : 0.5758  
  
Interpretation:  
→ R2 = 0.5758 means the model explains 57.6% of variance in price.  
→ Closer to 1.0 is better.  
  
=====  
PRACTICAL 4 COMPLETE  
=====  
  
(venv) shreehari@DESKTOP-0F0T52N: /mnt/d/Third Year/DSBDA/practicals$
```

PRACTICAL 5

➤ Source Code

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.preprocessing import StandardScaler

DIVIDER = "=" * 60

np.random.seed(42)
n = 100

age = np.random.randint(18, 60, n)
salary = np.random.randint(15000, 150000, n)
# Older + higher salary → more likely to buy
purchased = ((age > 35) & (salary > 70000)).astype(int)
# Add some noise
noise_idx = np.random.choice(n, 10, replace=False)
purchased[noise_idx] = 1 - purchased[noise_idx]

df = pd.DataFrame({'Age': age, 'EstimatedSalary': salary, 'Purchased': purchased})

print(f"\n{DIVIDER}")
print(" STEP 1 — Dataset Loaded")
print(f" Shape: {df.shape}")
print(DIVIDER)
print(df.head(10).to_string(index=False))
print(f"\n Purchased distribution:\n{df['Purchased'].value_counts()}")

X = df[['Age', 'EstimatedSalary']]
y = df['Purchased']

print(f"\n{DIVIDER}")
print(" STEP 2 — Features and Target")
print(DIVIDER)
print(f" X (features): Age, EstimatedSalary")
print(f" y (target) : Purchased (0=No, 1=Yes)")
```

```

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.25, random_state=42
)

print(f"\n{DIVIDER}")
print(" STEP 3 — Train-Test Split (75% train, 25% test)")
print(DIVIDER)
print(f" Training samples : {X_train.shape[0]}")
print(f" Testing samples : {X_test.shape[0]}")

scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

print(f"\n{DIVIDER}")
print(" STEP 4 — Feature Scaling done (StandardScaler)")
print(" (Makes Age and Salary comparable in scale)")
print(DIVIDER)

model = LogisticRegression()
model.fit(X_train, y_train)

print(f"\n{DIVIDER}")
print(" STEP 5 — Logistic Regression Model Trained!")
print(DIVIDER)
print(f" Coefficients: {model.coef_}")
print(f" Intercept : {model.intercept_}")
y_pred = model.predict(X_test)

print(f"\n{DIVIDER}")
print(" STEP 6 — Sample Predictions vs Actual")
print(DIVIDER)
result = pd.DataFrame({'Actual': y_test.values, 'Predicted': y_pred})
print(result.head(15).to_string(index=False))

cm = confusion_matrix(y_test, y_pred)

print(f"\n{DIVIDER}")
print(" STEP 7 — Confusion Matrix")
print(DIVIDER)

```

```

print(f"\n Confusion Matrix Layout:")
print(f"          Predicted 0 Predicted 1")
print(f" Actual 0 :   {cm[0,0]}       {cm[0,1]}")
print(f" Actual 1 :   {cm[1,0]}       {cm[1,1]}")
print(f"\n TN={cm[0,0]} FP={cm[0,1]} FN={cm[1,0]} TP={cm[1,1]}")

```

```

TP = cm[1, 1]
TN = cm[0, 0]
FP = cm[0, 1]
FN = cm[1, 0]

```

```

accuracy = (TP + TN) / (TP + TN + FP + FN)
error_rate = (FP + FN) / (TP + TN + FP + FN)
precision = TP / (TP + FP) if (TP + FP) != 0 else 0
recall = TP / (TP + FN) if (TP + FN) != 0 else 0
f1 = 2 * precision * recall / (precision + recall) if (precision + recall) != 0 else 0

```

```

print(f"\n{DIVIDER}")
print(" STEP 8 — Performance Metrics")
print(DIVIDER)
print(f" Accuracy : {accuracy:.4f} ({accuracy*100:.2f}%)")
print(f" Error Rate : {error_rate:.4f} ({error_rate*100:.2f}%)")
print(f" Precision : {precision:.4f}")
print(f" Recall : {recall:.4f}")
print(f" F1 Score : {f1:.4f}")

```

```

print(f"\n{DIVIDER}")
print(" STEP 9 — Full Classification Report (sklearn)")
print(DIVIDER)
print(classification_report(y_test, y_pred, target_names=['Not Purchased', 'Purchased']))

```

```

print(f"\n{'='*60}")
print(" PRACTICAL 5 COMPLETE")
print(f"\n{'='*60}")

```

➤ Output

```
Ubuntu
(venv) shreehari@DESKTOP-0F0T52N:/mnt/d/Third Year/DS&DA/practicals$ python3 prac_5.py

=====
STEP 1 - Dataset Loaded
Shape: (100, 3)
=====
Age  EstimatedSalary  Purchased
56      23392           0
46      45535           0
32     128569           0
25      67256           0
38     104135           1
56     142478           1
36      50222           0
40      92373           1
28     138684           0
28      25965           1

Purchased distribution:
Purchased
0      56
1      44
Name: count, dtype: int64

=====
STEP 2 - Features and Target
=====
X (features): Age, EstimatedSalary
y (target)  : Purchased (0=No, 1=Yes)

=====
STEP 3 - Train-Test Split (75% train, 25% test)
```

```
Ubuntu
Training samples : 75
Testing samples  : 25

=====
STEP 4 - Feature Scaling done (StandardScaler)
(Makes Age and Salary comparable in scale)
=====

STEP 5 - Logistic Regression Model Trained!
=====
Coefficients: [[1.22092479 0.69097723]]
Intercept   : [-0.24774001]

=====
STEP 6 - Sample Predictions vs Actual
=====
Actual Predicted
0      0
0      0
1      1
1      1
1      1
0      0
1      1
0      0
0      0
0      1
1      1
0      0
0      0
0      0
```

```
=====  
STEP 7 - Confusion Matrix  
=====
```

Confusion Matrix Layout:

	Predicted 0	Predicted 1
Actual 0 :	13	2
Actual 1 :	1	9

TN=13 FP=2 FN=1 TP=9

```
=====  
STEP 8 - Performance Metrics  
=====
```

Accuracy : 0.8800 (88.00%)
Error Rate : 0.1200 (12.00%)
Precision : 0.8182
Recall : 0.9000
F1 Score : 0.8571

```
=====  
STEP 9 - Full Classification Report (sklearn)  
=====
```

	precision	recall	f1-score	support
Not Purchased	0.93	0.87	0.90	15
Purchased	0.82	0.90	0.86	10
accuracy			0.88	25
macro avg	0.87	0.88	0.88	25
weighted avg	0.88	0.88	0.88	25

```
=====  
PRACTICAL 5 COMPLETE  
=====
```

(venv) shreehari@DESKTOP-0FOT52N: /mnt/0/Thlrd Year/BSBDA/practical:\$



PRACTICAL 6

➤ Source Code

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score

DIVIDER = "=" * 60

from sklearn.datasets import load_iris

iris_bunch = load_iris(as_frame=True)
df = iris_bunch.frame
df['species'] = df['target'].map({0: 'Iris-setosa', 1: 'Iris-versicolor', 2: 'Iris-virginica'})
df = df.drop('target', axis=1)

print(f"\n{DIVIDER}")
print(" STEP 1 — Iris Dataset Loaded")
print(f" Shape: {df.shape}")
print(DIVIDER)
print(df.head(10).to_string(index=False))

print(f"\n{DIVIDER}")
print(" STEP 2 — Dataset Info")
print(DIVIDER)
print(df.dtypes)
print(f"\n Species distribution:\n{df['species'].value_counts()}")

print(f"\n{DIVIDER}")
print(" STEP 3 — Statistical Summary")
print(DIVIDER)
print(df.describe().round(3))

X = df.drop('species', axis=1)
y = df['species']

print(f"\n{DIVIDER}")
print(" STEP 4 — Features (X) and Target (y)")
```



```

print(DIVIDER)
print(f" Features: {list(X.columns)}")
print(f" Target : species")

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42
)

print(f"\n{DIVIDER}")
print(" STEP 5 — Train-Test Split (70% train, 30% test)")
print(DIVIDER)
print(f" Training samples: {X_train.shape[0]}")
print(f" Testing samples: {X_test.shape[0]}")

model = GaussianNB()
model.fit(X_train, y_train)

print(f"\n{DIVIDER}")
print(" STEP 6 — Gaussian Naïve Bayes Model Trained!")
print(DIVIDER)
print(f" Classes learned: {list(model.classes_)}")

y_pred = model.predict(X_test)

print(f"\n{DIVIDER}")
print(" STEP 7 — Sample Predictions vs Actual")
print(DIVIDER)
result = pd.DataFrame({'Actual': y_test.values, 'Predicted': y_pred})
print(result.head(15).to_string(index=False))

cm = confusion_matrix(y_test, y_pred)
classes = ['Setosa', 'Versicolor', 'Virginica']

print(f"\n{DIVIDER}")
print(" STEP 8 — Confusion Matrix (3x3 — Multiclass)")
print(DIVIDER)
cm_df = pd.DataFrame(cm,
    index=[f'Actual {c}' for c in classes],
    columns=[f'Pred {c}' for c in classes])
print(cm_df)

```

```

print(f"\n{DIVIDER}")
print(" STEP 9 — Performance Metrics per Class")
print(DIVIDER)

for i, cls in enumerate(classes):
    TP = cm[i, i]
    FP = cm[:, i].sum() - TP
    FN = cm[i, :].sum() - TP
    TN = cm.sum() - TP - FP - FN

    accuracy = (TP + TN) / cm.sum()
    precision = TP / (TP + FP) if (TP + FP) != 0 else 0
    recall = TP / (TP + FN) if (TP + FN) != 0 else 0
    f1 = 2 * precision * recall / (precision + recall) if (precision + recall) != 0 else 0

    print(f"\n [{cls}]")
    print(f" TP={TP} FP={FP} FN={FN} TN={TN}")
    print(f" Accuracy : {accuracy:.4f}")
    print(f" Precision : {precision:.4f}")
    print(f" Recall : {recall:.4f}")
    print(f" F1 Score : {f1:.4f}")

overall = accuracy_score(y_test, y_pred)
print(f"\n{DIVIDER}")
print(" STEP 10 — Overall Accuracy & Full Report")
print(DIVIDER)
print(f" Overall Accuracy: {overall:.4f} ({overall*100:.2f}%)")
print(f"\n{classification_report(y_test, y_pred)}")

print(f"\n{'='*60}")
print(" PRACTICAL 6 COMPLETE")
print(f"\n{'='*60}")

```

➤ Output

```
Ubuntu
(venv) shreehari@DESKTOP-0F0T52N:/mnt/d/Third Year/DSBDA/practical:$ python3 prac_6.py

=====
STEP 1 - Iris Dataset Loaded
Shape: (150, 5)
=====
sepal length (cm)  sepal width (cm)  petal length (cm)  petal width (cm)  species
5.1                3.5                1.4                0.2 Iris-setosa
4.9                3.0                1.4                0.2 Iris-setosa
4.7                3.2                1.3                0.2 Iris-setosa
4.6                3.1                1.5                0.2 Iris-setosa
5.0                3.6                1.4                0.2 Iris-setosa
5.4                3.9                1.7                0.4 Iris-setosa
4.6                3.4                1.4                0.3 Iris-setosa
5.0                3.4                1.5                0.2 Iris-setosa
4.4                2.9                1.4                0.2 Iris-setosa
4.9                3.1                1.5                0.1 Iris-setosa

=====
STEP 2 - Dataset Info
=====
sepal length (cm)    float64
sepal width (cm)     float64
petal length (cm)    float64
petal width (cm)     float64
species              str
dtype: object

Species distribution:
species
Iris-setosa      50
Iris-versicolor  50

=====
```

```
Ubuntu
petal length (cm)    float64
petal width (cm)     float64
species              str
dtype: object

Species distribution:
species
Iris-setosa      50
Iris-versicolor  50
Iris-virginica   50
Name: count, dtype: int64

=====
STEP 3 - Statistical Summary
=====
count    sepal length (cm)  sepal width (cm)  petal length (cm)  petal width (cm)
mean      5.843             3.057             3.758             1.199
std       0.828             0.436             1.765             0.762
min       4.300             2.000             1.000             0.100
25%       5.100             2.800             1.600             0.300
50%       5.800             3.000             4.350             1.300
75%       6.400             3.300             5.100             1.800
max       7.900             4.400             6.900             2.500

=====
STEP 4 - Features (X) and Target (y)
=====
Features: ['sepal length (cm)', 'sepal width (cm)', 'petal length (cm)', 'petal width (cm)']
Target : species

=====
```

```
=====
STEP 5 - Train-Test Split (70% train, 30% test)
=====
Training samples: 105
Testing samples: 45

=====
STEP 6 - Gaussian Naïve Bayes Model Trained!
=====
Classes learned: [np.str_('Iris-setosa'), np.str_('Iris-versicolor'), np.str_('Iris-virginica')]

=====
STEP 7 - Sample Predictions vs Actual
=====
      Actual      Predicted
Iris-versicolor Iris-versicolor
Iris-setosa      Iris-setosa
Iris-virginica   Iris-virginica
Iris-versicolor Iris-versicolor
Iris-versicolor Iris-versicolor
Iris-setosa      Iris-setosa
Iris-versicolor Iris-versicolor
Iris-virginica   Iris-virginica
Iris-versicolor Iris-versicolor
Iris-versicolor Iris-versicolor
Iris-virginica   Iris-virginica
Iris-setosa      Iris-setosa
Iris-setosa      Iris-setosa
Iris-setosa      Iris-setosa
Iris-setosa      Iris-setosa

=====

=====
STEP 8 - Confusion Matrix (3x3 - Multiclass)
=====
      Pred Setosa  Pred Versicolor  Pred Virginica
Actual Setosa      19                0                0
Actual Versicolor  0                12               1
Actual Virginica   0                0                13

=====
STEP 9 - Performance Metrics per Class
=====

[Setosa]
TP=19 FP=0 FN=0 TN=26
Accuracy : 1.0000
Precision : 1.0000
Recall : 1.0000
F1 Score : 1.0000

[Versicolor]
TP=12 FP=0 FN=1 TN=32
Accuracy : 0.9778
Precision : 1.0000
Recall : 0.9231
F1 Score : 0.9600

[Virginica]
TP=13 FP=1 FN=0 TN=31
Accuracy : 0.9778
Precision : 0.9286
Recall : 1.0000
F1 Score : 0.9630
```

```
Ubuntu
Precision : 1.0000
Recall    : 0.9231
F1 Score  : 0.9600

[Virginica]
TP=13 FP=1 FN=0 TN=31
Accuracy  : 0.9778
Precision : 0.9286
Recall    : 1.0000
F1 Score  : 0.9630

=====
STEP 10 - Overall Accuracy & Full Report
=====
Overall Accuracy: 0.9778 (97.78%)

      precision    recall  f1-score   support

 Iris-setosa       1.00      1.00      1.00        19
 Iris-versicolor   1.00      0.92      0.96        13
 Iris-virginica     0.93      1.00      0.96        13

 accuracy
macro avg      0.98      0.97      0.97        45
weighted avg    0.98      0.98      0.98        45

=====
PRACTICAL 6 COMPLETE
=====

(env) shreehari@DESKTOP-0F0T52N: /mnt/d/Third Year/DSBDA/practicals$
```



PRACTICAL 7

➤ Source Code

```
import nltk
import pandas as pd
from sklearn.feature_extraction.text import TfidfVectorizer
print("Downloading NLTK data (only needed once)...")
nltk.download('punkt', quiet=True)
nltk.download('punkt_tab', quiet=True)
nltk.download('averaged_perceptron_tagger', quiet=True)
nltk.download('averaged_perceptron_tagger_eng', quiet=True)
nltk.download('stopwords', quiet=True)
nltk.download('wordnet', quiet=True)
from nltk.tokenize import word_tokenize
from nltk import pos_tag
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer, WordNetLemmatizer
DIVIDER = "=" * 60
text = "Data analytics is the process of analyzing raw data to find useful information."

docs = [
    "Data analytics is the process of analyzing raw data to find useful information.",
    "Machine learning algorithms learn from data and improve over time.",
    "Python is a popular programming language for data science and analytics.",
]

print(f"\n{DIVIDER}")
print(" STEP 0 — Sample Text")
print(DIVIDER)
print(f" Text: {text}")
tokens = word_tokenize(text)
print(f"\n{DIVIDER}")
print(" STEP 1 — Tokenization (splitting text into words)")
print(DIVIDER)
print(f" Tokens ({len(tokens)}): \n {tokens}")
pos_tags = pos_tag(tokens)
print(f"\n{DIVIDER}")
print(" STEP 2 — POS Tagging (Part of Speech)")
print(" NN=Noun VBZ=Verb JJ=Adjective IN=Preposition DT=Determiner")
print(DIVIDER)
```

```

for word, tag in pos_tags:
    print(f"  {word:15s} → {tag}")

stop_words = set(stopwords.words('english'))
filtered_tokens = [w for w in tokens if w.lower() not in stop_words and w.isalpha()]

print(f"\n{DIVIDER}")
print(" STEP 3 — Stop Words Removal")
print(f" Stop words removed (like: is, the, of, to, a ...)")
print(DIVIDER)
print(f" Before ({len(tokens)}) tokens: {tokens}")
print(f"\n After ({len(filtered_tokens)}) tokens: {filtered_tokens}")

ps = PorterStemmer()
stemmed = [ps.stem(w) for w in filtered_tokens]

print(f"\n{DIVIDER}")
print(" STEP 4 — Stemming (reduces to root, may not be real word)")
print(DIVIDER)
for orig, stem in zip(filtered_tokens, stemmed):
    print(f"  {orig:15s} → {stem}")
lm = WordNetLemmatizer()
lemmatized = [lm.lemmatize(w.lower()) for w in filtered_tokens]

print(f"\n{DIVIDER}")
print(" STEP 5 — Lemmatization (reduces to real base word)")
print(DIVIDER)
for orig, lemma in zip(filtered_tokens, lemmatized):
    print(f"  {orig:15s} → {lemma}")
print(f"\n{DIVIDER}")
print(" Comparison: Stemming vs Lemmatization")
print(DIVIDER)
comp = pd.DataFrame({
    'Original': filtered_tokens,
    'Stemmed': stemmed,
    'Lemmatized': lemmatized,
})
print(comp.to_string(index=False))

print(f"\n{DIVIDER}")
print(" STEP 6 — TF-IDF Representation")

```

```

print(" (Shows how important each word is in each document)")
print(DIVIDER)
print(" Documents used:")
for i, d in enumerate(docs):
    print(f" Doc{i+1}: {d}")

vectorizer = TfidfVectorizer()
tfidf_matrix = vectorizer.fit_transform(docs)

feature_names = vectorizer.get_feature_names_out()
tfidf_df = pd.DataFrame(
    tfidf_matrix.toarray(),
    columns=feature_names,
    index=[f'Doc{i+1}' for i in range(len(docs))]
)

print(f"\n TF-IDF Matrix (only top words shown):")
# Show top 8 words by total TF-IDF score
top_cols = tfidf_df.sum().nlargest(10).index
print(tfidf_df[top_cols].round(4))

print(f"\n Full vocabulary size: {len(feature_names)} words")
print(f" Matrix shape      : {tfidf_matrix.shape}")

print(f"\n{DIVIDER}")
print(" STEP 7 — Manual TF-IDF Calculation Example")
print(" Word: 'data' in Document 1")
print(DIVIDER)
word = 'data'
doc1_words = docs[0].lower().split()
tf = doc1_words.count(word) / len(doc1_words)
docs_with_word = sum(1 for d in docs if word in d.lower())
import math
idf = math.log(len(docs) / docs_with_word)
print(f" TF = {doc1_words.count(word)} / {len(doc1_words)} = {tf:.4f}")
print(f" IDF = log({len(docs)} / {docs_with_word}) = {idf:.4f}")
print(f" TF-IDF = TF × IDF = {tf:.4f} × {idf:.4f} = {tf*idf:.4f}")

print(f"\n{'='*60}")
print(" PRACTICAL 7 COMPLETE")
print(f"\n{'='*60}")

```


➤ Output

```
Ubuntu
(venv) shreehari@DESKTOP-0FO152N:/mnt/d/Third Year/DSBDA/practical:$ python3 prac_7.py
Downloading NLTK data (only needed once)...

=====
STEP 0 - Sample Text
=====
Text: Data analytics is the process of analyzing raw data to find useful information.
=====
STEP 1 - Tokenization (splitting text into words)
=====
Tokens (14):
['Data', 'analytics', 'is', 'the', 'process', 'of', 'analyzing', 'raw', 'data', 'to', 'find', 'useful', 'information', '.']
]

=====
STEP 2 - POS Tagging (Part of Speech)
NN=Noun VBZ=Verb JJ=Adjective IN=Preposition DT=Determiner
=====
Data      → NNP
analytics → NNS
is        → VBZ
the       → DT
process   → NN
of        → IN
analyzing → VBG
raw       → JJ
data      → NNS
to        → TO
find      → VB
useful    → JJ
information → NN

=====
```

```
Ubuntu
=====
STEP 3 - Stop Words Removal
Stop words removed (like: is, the, of, to, a ...)
=====
Before (14) tokens: ['Data', 'analytics', 'is', 'the', 'process', 'of', 'analyzing', 'raw', 'data', 'to', 'find', 'useful', 'information', '.']
After (9) tokens: ['Data', 'analytics', 'process', 'analyzing', 'raw', 'data', 'find', 'useful', 'information']

=====
STEP 4 - Stemming (reduces to root, may not be real word)
=====
Data      → data
analytics → analyt
process   → process
analyzing → analyz
raw       → raw
data      → data
find      → find
useful    → use
information → inform

=====
STEP 5 - Lemmatization (reduces to real base word)
=====
Data      → data
analytics → analytics
process   → process
analyzing → analyzing
raw       → raw
data      → data
find      → find

=====
```

```
=====  
Comparison: Stemming vs Lemmatization  
=====
```

Original	Stemmed	Lemmatized
Data	data	data
analytics	analyt	analytics
process	process	process
analyzing	analyz	analyzing
raw	raw	raw
data	data	data
find	find	find
useful	use	useful
information	inform	information

```
=====
```

STEP 6 – TF-IDF Representation
(Shows how important each word is in each document)

Documents used:
Doc1: Data analytics is the process of analyzing raw data to find useful information.
Doc2: Machine learning algorithms learn from data and improve over time.
Doc3: Python is a popular programming language for data science and analytics.

TF-IDF Matrix (only top words shown):

	data	and	analytics	is	for	language	popular	programming	python	science
Doc1	0.3475	0.0000	0.2238	0.2238	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000
Doc2	0.1977	0.2545	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000
Doc3	0.2077	0.2675	0.2675	0.2675	0.3517	0.3517	0.3517	0.3517	0.3517	0.3517

Full vocabulary size: 27 words
Matrix shape : (3, 27)

=====
STEP 7 – Manual TF-IDF Calculation Example
Word: 'data' in Document 1
=====

$TF = 2 / 13 = 0.1538$
 $IDF = \log(3 / 3) = 0.0000$
 $TF-IDF = TF \times IDF = 0.1538 \times 0.0000 = 0.0000$

=====
PRACTICAL 7 COMPLETE
=====

```
(venv) shreehari@DESKTOP-0FOT52N:/mnt/d/Third Year/DSBDA/practicals$
```

PRACTICAL 8

➤ Source Code

```
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd

DIVIDER = "=" * 60

df = sns.load_dataset('titanic')

print(f"\n{DIVIDER}")
print(" STEP 1 — Titanic Dataset Loaded (seaborn built-in)")
print(f" Shape: {df.shape}")
print(DIVIDER)
print(df.head())

print(f"\n{DIVIDER}")
print(" STEP 2 — Dataset Info & Missing Values")
print(DIVIDER)
print(df.isnull().sum())

print(f"\n{DIVIDER}")
print(" STEP 3 — Statistical Summary")
print(DIVIDER)
print(df.describe())

plt.figure(figsize=(6, 4))
sns.countplot(x='survived', data=df, palette='Set2')
plt.title("Plot 1: Survival Count (0=No, 1=Yes)")
plt.xlabel("Survived")
plt.ylabel("Count")
plt.tight_layout()
plt.savefig("plot1_survival_count.png")
plt.close()
print(f"\n [Saved] Plot 1 → plot1_survival_count.png")
print(" Shows: How many passengers survived vs did not survive.")

plt.figure(figsize=(6, 4))
sns.countplot(x='survived', hue='sex', data=df, palette='Set1')
```

```
plt.title("Plot 2: Survival by Gender")
plt.xlabel("Survived")
plt.ylabel("Count")
plt.legend(title='Gender')
plt.tight_layout()
plt.savefig("plot2_survival_by_gender.png")
plt.close()
print(f" [Saved] Plot 2 → plot2_survival_by_gender.png")
print(" Shows: Female passengers survived much more than males.")
```

```
plt.figure(figsize=(6, 4))
sns.countplot(x='survived', hue='pclass', data=df, palette='Set3')
plt.title("Plot 3: Survival by Passenger Class")
plt.xlabel("Survived")
plt.ylabel("Count")
plt.legend(title='Pclass')
plt.tight_layout()
plt.savefig("plot3_survival_by_class.png")
plt.close()
print(f" [Saved] Plot 3 → plot3_survival_by_class.png")
print(" Shows: 1st class passengers had higher survival rate.")
```

```
plt.figure(figsize=(7, 4))
sns.histplot(df['fare'].dropna(), bins=30, kde=True, color='steelblue')
plt.title("Plot 4: Fare Distribution")
plt.xlabel("Fare (ticket price)")
plt.ylabel("Frequency")
plt.tight_layout()
plt.savefig("plot4_fare_distribution.png")
plt.close()
print(f" [Saved] Plot 4 → plot4_fare_distribution.png")
print(" Shows: Most passengers paid low fares; few paid very high.")
```

```
plt.figure(figsize=(7, 4))
sns.histplot(df['age'].dropna(), bins=20, kde=True, color='coral')
plt.title("Plot 5: Age Distribution of Passengers")
plt.xlabel("Age")
plt.ylabel("Frequency")
plt.tight_layout()
plt.savefig("plot5_age_distribution.png")
plt.close()
```

```

print(f" [Saved] Plot 5 → plot5_age_distribution.png")
print(" Shows: Most passengers were between 20-40 years old.")

print(f"\n{DIVIDER}")
print(" OBSERVATIONS")
print(DIVIDER)
survived = df['survived'].value_counts()
gender_surv = df.groupby('sex')['survived'].mean()
class_surv = df.groupby('pclass')['survived'].mean()

print(f" Total Survived : {survived[1]} | Not Survived: {survived[0]}")
print(f"\n Survival Rate by Gender:")
for g, r in gender_surv.items():
    print(f" {g:8s}: {r*100:.1f}%")
print(f"\n Survival Rate by Class:")
for c, r in class_surv.items():
    print(f" Class {c}: {r*100:.1f}%")

print(f"\n{'='*60}")
print(" PRACTICAL 8 COMPLETE")
print(" 5 plots saved as PNG in current directory")
print(f"\n{'='*60}\n")

```

➤ Output

```
Ubuntu
(venv) shreehari@DESKTOP-0F0T52N:/mnt/d/Third Year/DSBDA/practical:$ python3 prac_8.py

=====
STEP 1 - Titanic Dataset Loaded (seaborn built-in)
Shape: (891, 15)
=====
survived  pclass  sex    age  sibsp  parch  fare  embarked  class  who  adult_male  deck  embark_town  alive  alone
0         0       3  male   22.0    1     0   7.2500     S  Third  man         True  NaN  Southampton  no    False
1         1       1 female   38.0    1     0  71.2833     C  First woman      False  C    Cherbourg   yes  False
2         1       3 female   26.0    0     0   7.9250     S  Third woman      False  NaN  Southampton  yes  True
3         1       1 female   35.0    1     0  53.1000     S  First woman      False  C    Southampton  yes  False
4         0       3  male   35.0    0     0   8.0500     S  Third  man         True  NaN  Southampton  no    True

=====
STEP 2 - Dataset Info & Missing Values
=====
survived      0
pclass        0
sex           0
age          177
sibsp         0
parch         0
fare          0
embarked      2
class         0
who           0
adult_male    0
deck         688
embark_town   2
alive         0
alone         0
dtype: int64
```

```
Ubuntu
=====
STEP 3 - Statistical Summary
=====
count  survived  pclass    age    sibsp    parch    fare
mean    0.383838   2.308642  29.699118  0.523008  0.381594  32.204208
std     0.486592   0.836071  14.526497  1.102743  0.806057  49.693429
min     0.000000   1.000000   0.420000  0.000000  0.000000  0.000000
25%     0.000000   2.000000  20.125000  0.000000  0.000000  7.910400
50%     0.000000   3.000000  28.000000  0.000000  0.000000  14.454200
75%     1.000000   3.000000  38.000000  1.000000  0.000000  31.000000
max     1.000000   3.000000  80.000000  8.000000  6.000000  512.329200
/mnt/d/Third Year/DSBDA/practicals/prac_8.py:26: FutureWarning:
Passing 'palette' without assigning 'hue' is deprecated and will be removed in v0.14.0. Assign the 'x' variable to 'hue' and
set 'legend=False' for the same effect.

sns.countplot(x='survived', data=df, palette='Set2')

[Saved] Plot 1 -> plot1_survival_count.png
Shows: How many passengers survived vs did not survive.
[Saved] Plot 2 -> plot2_survival_by_gender.png
Shows: Female passengers survived much more than males.
[Saved] Plot 3 -> plot3_survival_by_class.png
Shows: 1st class passengers had higher survival rate.
[Saved] Plot 4 -> plot4_fare_distribution.png
Shows: Most passengers paid low fares; few paid very high.
[Saved] Plot 5 -> plot5_age_distribution.png
Shows: Most passengers were between 20-40 years old.

=====
```

```
Ubuntu
[Saved] Plot 1 → plot1_survival_count.png
Shows: How many passengers survived vs did not survive.
[Saved] Plot 2 → plot2_survival_by_gender.png
Shows: Female passengers survived much more than males.
[Saved] Plot 3 → plot3_survival_by_class.png
Shows: 1st class passengers had higher survival rate.
[Saved] Plot 4 → plot4_fare_distribution.png
Shows: Most passengers paid low fares; few paid very high.
[Saved] Plot 5 → plot5_age_distribution.png
Shows: Most passengers were between 20-40 years old.


=====
OBSERVATIONS
=====
Total Survived : 342 | Not Survived: 549



Survival Rate by Gender:
female : 74.2%
male : 18.9%



Survival Rate by Class:
Class 1: 63.0%
Class 2: 47.3%
Class 3: 24.2%

=====
PRACTICAL 8 COMPLETE
5 plots saved as PNG in current directory
=====

(venv) shreehari@DESKTOP-0FDT52N: /mnt/d/Third Year/DSEDA/practicals$
```







PRACTICAL 9

➤ Source Code

```
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd

DIVIDER = "=" * 60

df = sns.load_dataset('titanic')

print(f"\n{DIVIDER}")
print(" STEP 1 — Dataset Loaded")
print(f" Shape: {df.shape}")
print(DIVIDER)
print(df[['sex', 'age', 'survived']].head(10))

print(f"\n{DIVIDER}")
print(" STEP 2 — Age Column Statistics")
print(DIVIDER)
print(f" Count : {df['age'].count()}")
print(f" Mean  : {df['age'].mean():.2f}")
print(f" Median : {df['age'].median():.2f}")
print(f" Std Dev: {df['age'].std():.2f}")
print(f" Min   : {df['age'].min()}")
print(f" Max   : {df['age'].max()}")
q1 = df['age'].quantile(0.25)
q3 = df['age'].quantile(0.75)
print(f" Q1 (25th percentile): {q1}")
print(f" Q3 (75th percentile): {q3}")
print(f" IQR           : {q3 - q1}")

print(f"\n{DIVIDER}")
print(" STEP 3 — Missing Values in Age")
print(DIVIDER)
print(f" Missing Age values: {df['age'].isna().sum()}")
print(" (seaborn's boxplot automatically skips NaN values)")

plt.figure(figsize=(8, 5))
```

```
sns.boxplot(x='sex', y='age', hue='survived', data=df, palette='Set2')
plt.title("Plot 1: Age Distribution by Gender and Survival")
plt.xlabel("Gender")
plt.ylabel("Age")
plt.legend(title='Survived', labels=['No (0)', 'Yes (1)'])
plt.tight_layout()
plt.savefig("plot6_boxplot_age_gender_survival.png")
plt.close()
print(f"\n [Saved] Plot 1 → plot6_boxplot_age_gender_survival.png")
```

```
plt.figure(figsize=(7, 5))
sns.boxplot(x='pclass', y='age', data=df, palette='pastel')
plt.title("Plot 2: Age Distribution by Passenger Class")
plt.xlabel("Passenger Class")
plt.ylabel("Age")
plt.tight_layout()
plt.savefig("plot7_boxplot_age_by_class.png")
plt.close()
print(f" [Saved] Plot 2 → plot7_boxplot_age_by_class.png")
```

```
plt.figure(figsize=(7, 5))
sns.boxplot(x='pclass', y='fare', data=df, palette='coolwarm')
plt.title("Plot 3: Fare Distribution by Passenger Class")
plt.xlabel("Passenger Class")
plt.ylabel("Fare")
plt.tight_layout()
plt.savefig("plot8_boxplot_fare_by_class.png")
plt.close()
print(f" [Saved] Plot 3 → plot8_boxplot_fare_by_class.png")
```

```
print(f"\n{DIVIDER}")
print(" STEP 4 — Age Statistics by Gender & Survival")
print(DIVIDER)
grouped = df.groupby(['sex', 'survived'])['age'].describe().round(2)
print(grouped)
```

```
print(f"\n{DIVIDER}")
print(" OBSERVATIONS (from Box Plots)")
print(DIVIDER)
print("""
1. GENDER & AGE:
```

- Male passengers have a wider age range than females.
- Median age for males is slightly higher (~30) than females (~28).

2. SURVIVAL & GENDER:

- Among survivors, females tend to be slightly younger.
- Non-surviving males have a higher median age.
- This supports the 'women and children first' evacuation pattern.

3. AGE & CLASS:

- 1st class passengers are generally older (median ~37).
- 3rd class passengers are younger (median ~25).

4. FARE & CLASS:

- 1st class fares are much higher and widely spread.
- 2nd and 3rd class fares are low and tightly clustered.

5. OUTLIERS:

- Some passengers aged above 70 appear as outliers.
- Very high fares in 1st class show as outliers.

""")

```
print(f"{'='*60}")
print(" PRACTICAL 9 COMPLETE")
print(" 3 plots saved as PNG in current directory")
print(f"{'='*60}\n")
```

➤ Output

```
Ubuntu
(venv) shreehari@DESKTOP-0F0T52N:/mnt/d/Third Year/DSBDA/practical:$ python3 prac_9.py

=====
STEP 1 - Dataset Loaded
Shape: (891, 15)
=====
   sex  age  survived
0  male  22.0         0
1  female 38.0         1
2  female 26.0         1
3  female 35.0         1
4  male  35.0         0
5  male  NaN         0
6  male  54.0         0
7  male   2.0         0
8  female 27.0         1
9  female 14.0         1

=====
STEP 2 - Age Column Statistics
=====
Count    : 714
Mean     : 29.70
Median   : 28.00
Std Dev  : 14.53
Min      : 0.42
Max      : 80.0
Q1 (25th percentile): 20.125
Q3 (75th percentile): 38.0
IQR      : 17.875

=====
```

```
Ubuntu
=====
STEP 3 - Missing Values in Age
=====
Missing Age values: 177
(seaborn's boxplot automatically skips NaN values)

[Saved] Plot 1 → plot6_boxplot_age_gender_survival.png
/mnt/d/Third Year/DSBDA/practicals/prac_9.py:48: FutureWarning:
Passing 'palette' without assigning 'hue' is deprecated and will be removed in v0.14.0. Assign the 'x' variable to 'hue' and
set 'legend=False' for the same effect.

sns.boxplot(x='pclass', y='age', data=df, palette='pastel')
[Saved] Plot 2 → plot7_boxplot_age_by_class.png
/mnt/d/Third Year/DSBDA/practicals/prac_9.py:58: FutureWarning:
Passing 'palette' without assigning 'hue' is deprecated and will be removed in v0.14.0. Assign the 'x' variable to 'hue' and
set 'legend=False' for the same effect.

sns.boxplot(x='pclass', y='fare', data=df, palette='coolwarm')
[Saved] Plot 3 → plot8_boxplot_fare_by_class.png

=====
STEP 4 - Age Statistics by Gender & Survival
=====
   sex  survived  count  mean  std  min  25%  50%  75%  max
female 0         64.0  25.05  13.62  2.00  16.75  24.5  33.25  57.0
       1        197.0  28.85  14.18  0.75  19.00  28.0  38.00  63.0
male   0        360.0  31.62  14.06  1.00  21.75  29.0  39.25  74.0
       1         93.0  27.28  16.50  0.42  18.00  28.0  36.00  80.0
```

```
=====
OBSERVATIONS (from Box Plots)
=====

1. GENDER & AGE:
  - Male passengers have a wider age range than females.
  - Median age for males is slightly higher (~30) than females (~28).

2. SURVIVAL & GENDER:
  - Among survivors, females tend to be slightly younger.
  - Non-surviving males have a higher median age.
  - This supports the 'women and children first' evacuation pattern.


3. AGE & CLASS:
  - 1st class passengers are generally older (median ~37).
  - 3rd class passengers are younger (median ~25).



4. FARE & CLASS:
  - 1st class fares are much higher and widely spread.
  - 2nd and 3rd class fares are low and tightly clustered.

5. OUTLIERS:
  - Some passengers aged above 70 appear as outliers.
  - Very high fares in 1st class show as outliers.

=====
PRACTICAL 9 COMPLETE
3 plots saved as PNG in current directory
=====

(venv) shreehari@DESKTOP-0FDT52N:/mnt/d/Third Year/DSBDA/practicals$
```





PRACTICAL 10

➤ Source Code

```
from collections import defaultdict
DIVIDER = "=" * 60
input_text = ""

hadoop is simple
hadoop is powerful
map reduce in hadoop

data processing with hadoop is fast
map reduce is a programming model
"""

print(f"\n{DIVIDER}")
print(" PRACTICAL 11 — WordCount using MapReduce Logic")
print(DIVIDER)
print(" Input Text:")
print(input_text)

def mapper(text):
    """Tokenizes text and emits (word, 1) pairs."""
    intermediate = []
    for line in text.strip().split('\n'):
        for word in line.strip().split():
            word = word.lower()
            intermediate.append((word, 1))
    return intermediate

mapper_output = mapper(input_text)

print(f"{DIVIDER}")
print(" MAPPER OUTPUT — (word, 1) pairs emitted")
print(DIVIDER)
for pair in mapper_output:
    print(f" {pair[0]:15s} → {pair[1]}")

def shuffle_and_sort(mapper_output):
    """Groups values by key."""
```



```

grouped = defaultdict(list)
for key, value in mapper_output:
    grouped[key].append(value)
return dict(sorted(grouped.items()))

shuffled = shuffle_and_sort(mapper_output)

print(f"\n{DIVIDER}")
print(" SHUFFLE & SORT — Values grouped by key (word)")

print(DIVIDER)
for word, values in shuffled.items():
    print(f" {word:15s} → {values}")

def reducer(shuffled):
    """Sums values for each key."""
    result = {}
    for key, values in shuffled.items():
        result[key] = sum(values)
    return result
final_output = reducer(shuffled)
print(f"\n{DIVIDER}")

print(" REDUCER OUTPUT — Final Word Count")
print(DIVIDER)
print(f" {'Word':<20} Count")
print(f" {'-'*30}")

for word, count in sorted(final_output.items(), key=lambda x: -x[1]):
    print(f" {word:<20} {count}")

print(f"\n{DIVIDER}")

print(" SUMMARY")
print(DIVIDER)
print(f" Total unique words : {len(final_output)}")
print(f" Total word tokens : {sum(final_output.values())}")
print(f" Most frequent word : '{max(final_output, key=final_output.get)}'")

```

```
print(f"\n{DIVIDER}")
print(" MapReduce Flow (how Hadoop works)")
print(DIVIDER)
print("""
INPUT FILE
    ↓
[MAPPER] — reads lines, emits (word, 1) for each word
    ↓
[SHUFFLE & SORT] — groups all (word, [1,1,1,...]) together
    ↓
[REDUCER] — sums the list → emits (word, total_count)
    ↓
OUTPUT FILE
""")

print(f"{'='*60}")
print(" PRACTICAL 11 COMPLETE")
print(f"{'='*60}\n")
```

➤ Output

```
Ubuntu
x + v
(venv) shreehari@DESKTOP-0FDT52N:/mnt/d/Third Year/D580A/practicals$ python3 prac_10.py

=====
STEP 1 - Iris Dataset Loaded (seaborn built-in)
Shape: (150, 5)
=====
   sepal_length  sepal_width  petal_length  petal_width  species
0             5.1           3.5           1.4           0.2  setosa
1             4.9           3.0           1.4           0.2  setosa
2             4.7           3.2           1.3           0.2  setosa
3             4.6           3.1           1.5           0.2  setosa
4             5.0           3.6           1.4           0.2  setosa
5             5.4           3.9           1.7           0.4  setosa
6             4.6           3.4           1.4           0.3  setosa
7             5.0           3.4           1.5           0.2  setosa
8             4.4           2.9           1.4           0.2  setosa
9             4.9           3.1           1.5           0.1  setosa

=====
STEP 2 - Data Types of Each Column
=====
sepal_length    float64
sepal_width     float64
petal_length     float64
petal_width     float64
species         str
dtype: object

Numeric features : sepal_length, sepal_width, petal_length, petal_width
Categorical (nominal): species

Species distribution:
species
setosa      50
versicolor 50
virginica   50
Name: count, dtype: int64
```

```
Ubuntu
x + v

=====
STEP 3 - Statistical Summary
=====
   count  sepal_length  sepal_width  petal_length  petal_width
count    150.000      150.000      150.000      150.000
mean      5.013        3.057        3.758        1.199
std       0.828        0.436        1.765        0.762
min       4.300        2.000        1.000        0.100
25%       5.100        2.800        1.600        0.300
50%       5.000        3.000        3.300        1.300
75%       6.400        3.300        4.300        1.800
max       7.900        4.400        6.900        2.500

=====
STEP 4 - Missing Values Check
=====
sepal_length    0
sepal_width     0
petal_length     0
petal_width     0
species         0
dtype: int64
(Iris dataset is clean - no missing values!)

[Saved] Plot 1 -> plot9_iris_histograms.png
[Saved] Plot 2 -> plot10_iris_boxplot_all.png
/mnt/d/Third Year/D580A/practicals/prac_10.py:65: FutureWarning:
Passing 'palette' without assigning 'hue' is deprecated and will be removed in v0.14.0. Assign the 'x' variable to 'hue' and set 'legend=False' for the same effect.

sns.boxplot(x='species', y=col, data=df, palette='pastel', ax=ax)
/mnt/d/Third Year/D580A/practicals/prac_10.py:65: FutureWarning:
Passing 'palette' without assigning 'hue' is deprecated and will be removed in v0.14.0. Assign the 'x' variable to 'hue' and set 'legend=False' for the same effect.
```

```
=====
STEP 5 - Per-Feature Statistical Analysis
=====

[SEPAL LENGTH]
Mean : 5.843
Median : 5.888
Std Dev : 0.828
Q1=5.10 Q3=6.40 IQR=1.30
Outlier bounds: [3.15, 8.35]
Outliers found: 0

[SEPAL WIDTH]
Mean : 3.057
Median : 3.088
Std Dev : 0.436
Q1=2.80 Q3=3.30 IQR=0.50
Outlier bounds: [2.05, 4.05]
Outliers found: 4

[PETAL LENGTH]
Mean : 3.758
Median : 4.358
Std Dev : 1.765
Q1=1.60 Q3=5.10 IQR=3.50
Outlier bounds: [-3.65, 10.35]
Outliers found: 0

[PETAL WIDTH]
Mean : 1.199
Median : 1.388
Std Dev : 0.762
Q1=0.30 Q3=1.80 IQR=1.50
Outlier bounds: [-1.95, 4.05]
Outliers found: 0

=====
```

```
=====
STEP 6 - Mean of Features per Species
=====
      sepal_length sepal_width petal_length petal_width
species
setosa           5.006      3.428         1.462         0.246
versicolor       5.936      2.770         4.260         1.326
virginica         6.588      2.974         5.552         2.026

=====
OBSERVATIONS
=====

1. SEPAL LENGTH:
  - Moderate spread across all species.
  - Virginica has the largest sepal length.

2. SEPAL WIDTH:
  - Narrow range overall.
  - Setosa has widest sepal width.
  - Contains a few notable outliers.



3. PETAL LENGTH:
  - Clear separation between species (great for classification!).
  - Setosa: 1.0-2.0 cm | Versicolor: 3.0-5.0 cm | Virginica: 5.0-7.0 cm.


4. PETAL WIDTH:
  - Similar pattern to petal length.
  - Very strong feature for separating species.

KEY INSIGHT:
  Petal features (length & width) are more useful for
  classification than sepal features because they show
  clear separation between the 3 species.

=====
PRACTICAL 10 COMPLETE
3 plots saved as PNG in current directory
=====

(venv) shreeharl@DESKTOP-0F0T52N:/mnt/d/Third Year/DSBA/practicals$
```





Practical No. 11

WordCount.java:

```
import java.io.IOException;

import java.util.StringTokenizer;

import org.apache.hadoop.conf.Configuration;

import org.apache.hadoop.fs.Path;

import org.apache.hadoop.io.IntWritable;

import org.apache.hadoop.io.Text;

import org.apache.hadoop.mapreduce.Job;

import org.apache.hadoop.mapreduce.Mapper;

import org.apache.hadoop.mapreduce.Reducer;

import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;

import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class WordCount {

    public static class TokenizerMapper

        extends Mapper<Object, Text, Text, IntWritable> {

        private final static IntWritable one = new IntWritable(1);

        private Text word = new Text();

        public void map(Object key, Text value, Context context)

            throws IOException, InterruptedException {

            StringTokenizer itr = new StringTokenizer(value.toString());

            while (itr.hasMoreTokens()) {

                word.set(itr.nextToken().toLowerCase());

                context.write(word, one);

            }

        }

    }

    public static class IntSumReducer

        extends Reducer<Text, IntWritable, Text, IntWritable> {

        private IntWritable result = new IntWritable();

        public void reduce(Text key, Iterable<IntWritable> values,

            Context context)

            throws IOException, InterruptedException {

            int sum = 0;

            for (IntWritable val : values) {

                sum += val.get();

            }

        }

    }

}
```

```

        result.set(sum);

        context.write(key, result);
    }
}

public static void main(String[] args)

    throws Exception {
    if (args.length != 2) {

        System.err.println("Usage: WordCount <input path> <output path>");

        System.exit(-1);
    }

    Configuration conf = new Configuration();

    Job job = Job.getInstance(conf, "Word Count");

    job.setJarByClass(WordCount.class);

    job.setMapperClass(TokenizerMapper.class);

    job.setReducerClass(IntSumReducer.class);

    job.setOutputKeyClass(Text.class);

    job.setOutputValueClass(IntWritable.class);

    FileInputFormat.addInputPath(job, new Path(args[0]));

    FileOutputFormat.setOutputPath(job, new Path(args[1]));

    System.exit(job.waitForCompletion(true) ? 0 : 1);

}
}

```

sample.txt:

```

hello world

hello hadoop

hello world

```

part-r-00000.:

```

hadoop 1

hello 3

world 2

```


OUTPUT:

```
Activities Terminal Apr 21 11:57
lab5@lab5-Latitude-7490: ~/Desktop/Hadoop_Practical

lab5@lab5-Latitude-7490:~/Desktop/Hadoop_Practical$ javac -classpath $(hadoop classpath) -d wordcount_classes WordCount.java
javac: directory not found: wordcount_classes
Usage: javac <options> <source files>
use -help for a list of possible options
lab5@lab5-Latitude-7490:~/Desktop/Hadoop_Practical$ ls
input output jar WordCount.java
lab5@lab5-Latitude-7490:~/Desktop/Hadoop_Practical$ ls
input output wordcount.jar WordCount.java
lab5@lab5-Latitude-7490:~/Desktop/Hadoop_Practical$ javac -classpath $(hadoop classpath) -d wordcount_classes WordCount.java
javac: directory not found: wordcount_classes
Usage: javac <options> <source files>
use -help for a list of possible options
lab5@lab5-Latitude-7490:~/Desktop/Hadoop_Practical$ ls
input output wordcount_classes wordcount.jar WordCount.java
lab5@lab5-Latitude-7490:~/Desktop/Hadoop_Practical$ javac -classpath $(hadoop classpath) -d wordcount_classes WordCount.java
lab5@lab5-Latitude-7490:~/Desktop/Hadoop_Practical$ jar -cvf wordcount.jar -C wordcount_classes/ .
added manifest
adding: WordCount.class(in = 1617) (out= 898)(deflated 44%)
adding: WordCount$IntSumReducer.class(in = 1739) (out= 742)(deflated 57%)
adding: WordCount$TokenizerMapper.class(in = 1705) (out= 777)(deflated 56%)
```

```
lab5@lab5-Latitude-7490:~/Desktop/Hadoop_Practical$ hadoop jar wordcount.jar WordCount input output
2026-04-21 11:55:45,263 INFO impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2026-04-21 11:55:45,326 INFO impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2026-04-21 11:55:45,326 INFO impl.MetricsSystemImpl: JobTracker metrics system started
2026-04-21 11:55:45,407 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
2026-04-21 11:55:45,493 INFO input.FileInputFormat: Total input files to process : 1
2026-04-21 11:55:45,514 INFO mapreduce.JobSubmitter: number of splits:1
2026-04-21 11:55:45,653 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local466340237_0001
2026-04-21 11:55:45,653 INFO mapreduce.JobSubmitter: Executing with tokens: []
2026-04-21 11:55:45,790 INFO mapreduce.Job: The url to track the jobs is http://localhost:8080/
2026-04-21 11:55:45,790 INFO mapreduce.Job: Running job: job_local466340237_0001
2026-04-21 11:55:45,790 INFO mapreduce.LocalJobRunner: OutputCommitter set in config null
2026-04-21 11:55:45,795 INFO output.PathOutputCommitterFactory: No output committer factory defined, defaulting to FileOutputCommitterFactory
2026-04-21 11:55:45,796 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2026-04-21 11:55:45,796 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup_temporary folders under output directory:false, ignore cleanup failures: false
2026-04-21 11:55:45,796 INFO mapreduce.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapreduce.lib.output.FileOutputCommitter
2026-04-21 11:55:45,830 INFO mapreduce.LocalJobRunner: Waiting for map tasks
2026-04-21 11:55:45,830 INFO mapreduce.LocalJobRunner: Starting task: attempt_local466340237_0001_m_000000_0
2026-04-21 11:55:45,849 INFO output.PathOutputCommitterFactory: No output committer factory defined, defaulting to FileOutputCommitterFactory
2026-04-21 11:55:45,849 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2026-04-21 11:55:45,849 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup_temporary folders under output directory:false, ignore cleanup failures: false
2026-04-21 11:55:45,871 INFO mapreduce.Task: Using ResourceCalculatorProcessTree: []
2026-04-21 11:55:45,875 INFO mapreduce.MapTask: Processing split: file:/home/lab5/Desktop/Hadoop_Practical/input/sample.txt:0+38
2026-04-21 11:55:45,942 INFO mapreduce.MapTask: (EQUATOR) 0 kvl 26214396(104857584)
2026-04-21 11:55:45,942 INFO mapreduce.MapTask: mapreduce.task.io.sort.mb: 100
2026-04-21 11:55:45,942 INFO mapreduce.MapTask: soft limit at 83886080
2026-04-21 11:55:45,942 INFO mapreduce.MapTask: bufstart = 0; bufend = 104857000
2026-04-21 11:55:45,942 INFO mapreduce.MapTask: kvstart = 26214396; length = 6553600
2026-04-21 11:55:45,945 INFO mapreduce.MapTask: Map output collector class = org.apache.hadoop.mapreduce.MapTask$MapOutputBuffer
2026-04-21 11:55:45,952 INFO mapreduce.LocalJobRunner:
2026-04-21 11:55:45,953 INFO mapreduce.MapTask: Starting flush of map output
2026-04-21 11:55:45,953 INFO mapreduce.MapTask: Spilling map output
2026-04-21 11:55:45,953 INFO mapreduce.MapTask: bufstart = 0; bufend = 01; bufvoid = 104857000
2026-04-21 11:55:45,953 INFO mapreduce.MapTask: kvstart = 26214396(104857584); kvend = 26214376(104857584); length = 21/6553600
2026-04-21 11:55:45,958 INFO mapreduce.MapTask: Finished spill 0
2026-04-21 11:55:45,968 INFO mapreduce.Task: Task:attempt_local466340237_0001_m_000000_0 is done. And is in the process of committing
2026-04-21 11:55:45,970 INFO mapreduce.LocalJobRunner: map
2026-04-21 11:55:45,970 INFO mapreduce.Task: Task 'attempt_local466340237_0001_m_000000_0' done.
2026-04-21 11:55:45,970 INFO mapreduce.Task: Final counters for attempt_local466340237_0001_m_000000_0: Counters: 17
File System Counters
FILE: Number of bytes read=3393
FILE: Number of bytes written=635495
FILE: Number of read operations=0
```

```
Activities Terminal Apr 21 11:57
lab5@lab5-Latitude-7490: ~/Desktop/Hadoop_Practical

WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Output Format Counters
Bytes Written=37
2026-04-21 11:55:46,074 INFO mapreduce.LocalJobRunner: Finishing task: attempt_local466340237_0001_r_000000_0
2026-04-21 11:55:46,074 INFO mapreduce.LocalJobRunner: reduce task executor complete.
2026-04-21 11:55:46,795 INFO mapreduce.Job: Job job_local466340237_0001 running in uber mode : false
2026-04-21 11:55:46,797 INFO mapreduce.Job: map 100% reduce 100%
2026-04-21 11:55:46,799 INFO mapreduce.Job: Job job_local466340237_0001 completed successfully
2026-04-21 11:55:46,816 INFO mapreduce.Job: Counters: 30
File System Counters
FILE: Number of bytes read=6976
FILE: Number of bytes written=1271100
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
Map-Reduce Framework
Map input records=4
Map output records=6
Map output bytes=61
Map output materialized bytes=79
Input split bytes=102
Combine input records=0
Combine output records=0
Reduce input groups=3
Reduce shuffle bytes=79
Reduce input records=6
Reduce output records=3
Spilled Records=12
Shuffled Maps=1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=6
Total committed heap usage (bytes)=499122176
Shuffle
Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
Bytes Read=38
File Output Format Counters
Bytes Written=37
lab5@lab5-Latitude-7490:~/Desktop/Hadoop_Practical$ cat output/part-r-00000
hadoop 1
hello 3
world 2
lab5@lab5-Latitude-7490:~/Desktop/Hadoop_Practical$
```

Practical No. 12

LogAnalysis.java:

```
import java.io.IOException;

import org.apache.hadoop.conf.Configuration;

import org.apache.hadoop.fs.Path;

import org.apache.hadoop.io.IntWritable;

import org.apache.hadoop.io.Text;

import org.apache.hadoop.mapreduce.Job;

import org.apache.hadoop.mapreduce.Mapper;

import org.apache.hadoop.mapreduce.Reducer;

import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;

import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class LogAnalysis {

    public static class LogMapper

        extends Mapper<Object, Text, Text, IntWritable> {

            private final static IntWritable one = new IntWritable(1);

            private Text word = new Text();

            public void map(Object key, Text value, Context context)

                throws IOException, InterruptedException {

                String line = value.toString();

                if(line.contains("INFO")) {

                    word.set("INFO");

                    context.write(word, one);

                }

                if(line.contains("ERROR")) {

                    word.set("ERROR");

                    context.write(word, one);

                }

                if(line.contains("WARN")) {

                    word.set("WARN");

                    context.write(word, one);

                }

            }

        }

    public static class LogReducer

        extends Reducer<Text, IntWritable, Text, IntWritable> {
```

```

public void reduce(Text key, Iterable<IntWritable> values,
                    Context context)
    throws IOException, InterruptedException {
    int sum = 0;
    for(IntWritable val : values) {
        sum += val.get();
    }
    context.write(key, new IntWritable(sum));
}

}

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "log analysis");
    job.setJarByClass(LogAnalysis.class);
    job.setMapperClass(LogMapper.class);
    job.setReducerClass(LogReducer.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}

}

```

logs.txt:

```

2026 INFO User Login
2026 ERROR Database Failed
2026 INFO File Opened
2026 WARN Memory Low
2026 ERROR Network Failed
2026 INFO Logout

```

OUTPUT:

```
lab5@lab5-Latitude-7490:~/Desktop/Hadoop_12$ mkdir classes
mkdir: cannot create directory 'classes': File exists
lab5@lab5-Latitude-7490:~/Desktop/Hadoop_12$ javac -classpath $HADOOP_HOME/share/hadoop/common/*:$HADOOP_HOME/share/hadoop/mapreduce/* \
  -srcsrc/LogAnalysis.java -d classes
lab5@lab5-Latitude-7490:~/Desktop/Hadoop_12$ jar -cvf loganalysis.jar -C classes /
added manifest
adding: LogAnalysis$LogMapper.class (ln = 1784) (out= 758)(deflated 57%)
adding: LogAnalysis.class (ln = 1618) (out= 895)(deflated 44%)
adding: LogAnalysis$LogReducer.class (ln = 1641) (out= 888)(deflated 58%)
lab5@lab5-Latitude-7490:~/Desktop/Hadoop_12$ start-dfs.sh
Starting namenodes on [lab5-Latitude-7490]
lab5-Latitude-7490: lab5@lab5-latitude-7490: Permission denied (publickey,password).
Starting datanodes
localhost: lab5@localhost: Permission denied (publickey,password).
Starting secondary namenodes [lab5-Latitude-7490]
lab5-Latitude-7490: lab5@lab5-latitude-7490: Permission denied (publickey,password).
lab5@lab5-Latitude-7490:~/Desktop/Hadoop_12$ start-yarn.sh
Starting resource manager
resource manager is running as process 3883. Stop it first and ensure /tmp/hadoop-lab5-resource-manager.pid file is empty before retry.
Starting node managers
localhost: lab5@localhost: Permission denied (publickey,password).
```

```
lab5@lab5-Latitude-7490:~/Desktop/Hadoop_12$ ls
classes input LogAnalysis.jar output src
lab5@lab5-Latitude-7490:~/Desktop/Hadoop_12$ jps
11872 ResourceManager
11592 DataNode
12965 Jps
12028 NodeManager
11228 NameNode
11602 SecondaryNameNode
lab5@lab5-Latitude-7490:~/Desktop/Hadoop_12$ javac -classpath 'hadoop classpath' -d classes src/LogAnalysis.java
lab5@lab5-Latitude-7490:~/Desktop/Hadoop_12$ hdfs dfs -mkdir /input
lab5@lab5-Latitude-7490:~/Desktop/Hadoop_12$ hdfs dfs -put input/log.txt /input
put: 'input/log.txt': No such file or directory
lab5@lab5-Latitude-7490:~/Desktop/Hadoop_12$ hdfs dfs -put input/logs.txt /inputlab5@lab5-Latitude-7490:~/Desktop/Hadoop_12$ hdfs dfs -ls /input
Found 1 items
-rw-r--r-- 1 lab5 supergroup 134 2026-04-21 13:58 /input/logs.txt
lab5@lab5-Latitude-7490:~/Desktop/Hadoop_12$ hadoop jar LogAnalysis.jar LogAnalysis /input /output
2026-04-21 13:59:06,695 INFO Impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2026-04-21 13:59:06,802 INFO Impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2026-04-21 13:59:06,802 INFO Impl.MetricsSystemImpl: JobTracker metrics system started
2026-04-21 13:59:06,972 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
2026-04-21 13:59:07,126 INFO Input.FileInputFormat: Total input files to process : 1
2026-04-21 13:59:07,152 INFO mapreduce.JobSubmitter: number of splits:1
2026-04-21 13:59:07,283 INFO mapreduce.JobSubmitter: Submitting tokens for Job: job_local1633189946_0001
2026-04-21 13:59:07,283 INFO mapreduce.JobSubmitter: Executing with tokens: []
2026-04-21 13:59:07,412 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
2026-04-21 13:59:07,413 INFO mapreduce.Job: Running Job: job_local1633189946_0001
2026-04-21 13:59:07,414 INFO mapred.LocalJobRunner: OutputCommitter set in config null
2026-04-21 13:59:07,422 INFO output.PathOutputCommitterFactory: No output committer factory defined, defaulting to FileOutputCommitterFactory
2026-04-21 13:59:07,423 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2026-04-21 13:59:07,423 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false
2026-04-21 13:59:07,424 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapreduce.lib.output.FileOutputCommitter
2026-04-21 13:59:07,407 INFO mapred.LocalJobRunner: Waiting for map tasks
2026-04-21 13:59:07,468 INFO mapred.LocalJobRunner: Starting task: attempt_local1633189946_0001_m_000000_0
2026-04-21 13:59:07,493 INFO output.PathOutputCommitterFactory: No output committer factory defined, defaulting to FileOutputCommitterFactory
2026-04-21 13:59:07,493 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2026-04-21 13:59:07,493 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false
2026-04-21 13:59:07,507 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
2026-04-21 13:59:07,512 INFO mapred.MapTask: Processing split: hdfs://localhost:9000/input/logs.txt:0+134
2026-04-21 13:59:07,542 INFO mapred.MapTask: (EQUATOR) 0 kvt 26214396(104857584)
2026-04-21 13:59:07,543 INFO mapred.MapTask: mapreduce.task.io.sort.mb: 100
2026-04-21 13:59:07,543 INFO mapred.MapTask: soft limit at 83886080
2026-04-21 13:59:07,543 INFO mapred.MapTask: bufstart = 0; bufvoid = 104857600
2026-04-21 13:59:07,543 INFO mapred.MapTask: kvstart = 26214396; length = 6553600
2026-04-21 13:59:07,546 INFO mapred.MapTask: Map output collector class = org.apache.hadoop.mapred.MapTask$MapOutputBuffer
2026-04-21 13:59:07,806 INFO mapred.LocalJobRunner:
2026-04-21 13:59:07,806 INFO mapred.MapTask: Starting flush of map output
2026-04-21 13:59:07,808 INFO mapred.MapTask: Spilling map output
2026-04-21 13:59:07,808 INFO mapred.MapTask: bufstart = 0; bufend = 56; bufvoid = 104857600
2026-04-21 13:59:07,808 INFO mapred.MapTask: kvstart = 26214396(104857584); kvend = 26214376(104857584); length = 21/6553600
2026-04-21 13:59:07,815 INFO mapred.MapTask: Finished spill 0
2026-04-21 13:59:07,828 INFO mapred.Task: Task:attempt_local1633189946_0001_m_000000_0 is done. And is in the process of committing
2026-04-21 13:59:07,833 INFO mapred.LocalJobRunner: map
2026-04-21 13:59:07,833 INFO mapred.Task: Task 'attempt_local1633189946_0001_m_000000_0' done.
2026-04-21 13:59:07,840 INFO mapred.Task: Final counters for attempt_local1633189946_0001_m_000000_0: Counters: 23
File System Counters
FILE: Number of bytes read=134
FILE: Number of bytes written=630958
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=134
HDFS: Number of bytes written=0
HDFS: Number of read operations=5
HDFS: Number of large read operations=0
HDFS: Number of write operations=1
HDFS: Number of bytes read erasure-coded=0
Map-Reduce Framework
Map input records=6
Map output records=6
Map output bytes=56
Map output materialized bytes=74
Input split bytes=101
Combine input records=0
Spilled Records=6
Failed Shuffles=0
Merged Map outputs=0
GC time elapsed (ms)=112
Total committed heap usage (bytes)=549453824
File Input Format Counters
Bytes Read=134
2026-04-21 13:59:07,840 INFO mapred.LocalJobRunner: Finishing task: attempt_local1633189946_0001_m_000000_0
2026-04-21 13:59:07,841 INFO mapred.LocalJobRunner: Map task executor complete.
2026-04-21 13:59:07,844 INFO mapred.LocalJobRunner: Waiting for reduce tasks
2026-04-21 13:59:07,844 INFO mapred.LocalJobRunner: Starting task: attempt_local1633189946_0001_r_000000_0
2026-04-21 13:59:07,853 INFO output.PathOutputCommitterFactory: No output committer factory defined, defaulting to FileOutputCommitterFactory
2026-04-21 13:59:07,853 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2026-04-21 13:59:07,853 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false
2026-04-21 13:59:07,854 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
2026-04-21 13:59:07,857 INFO reduce.ShuffleTask: Using ShuffleConsumerPlugin: org.apache.hadoop.mapreduce.task.reduce.Shuffle@71479053
2026-04-21 13:59:07,859 WARN Impl.MetricsSystemImpl: JobTracker metrics system already initialized!
2026-04-21 13:59:07,879 INFO reduce.MergeManagerImpl: MergeManager: memoryLimit=2589196288, maxSingleShuffleLimit=647299072, mergeThreshold=1708809632, ioSortFactor=10, memToMemMergeOutputsThreshold=10
2026-04-21 13:59:07,882 INFO reduce.EventFetcher: attempt_local1633189946_0001_r_000000_0 Thread started: EventFetcher for fetching Map Completion Events
```

```
2026-04-21 13:59:07,543 INFO mapred.MapTask: bufstart = 0; bufvoid = 104857600
2026-04-21 13:59:07,543 INFO mapred.MapTask: kvstart = 26214396; length = 6553600
2026-04-21 13:59:07,546 INFO mapred.MapTask: Map output collector class = org.apache.hadoop.mapred.MapTask$MapOutputBuffer
2026-04-21 13:59:07,806 INFO mapred.LocalJobRunner:
2026-04-21 13:59:07,808 INFO mapred.MapTask: Starting flush of map output
2026-04-21 13:59:07,808 INFO mapred.MapTask: Spilling map output
2026-04-21 13:59:07,808 INFO mapred.MapTask: bufstart = 0; bufend = 56; bufvoid = 104857600
2026-04-21 13:59:07,808 INFO mapred.MapTask: kvstart = 26214396(104857584); kvend = 26214376(104857584); length = 21/6553600
2026-04-21 13:59:07,815 INFO mapred.MapTask: Finished spill 0
2026-04-21 13:59:07,828 INFO mapred.Task: Task:attempt_local1633189946_0001_m_000000_0 is done. And is in the process of committing
2026-04-21 13:59:07,833 INFO mapred.LocalJobRunner: map
2026-04-21 13:59:07,833 INFO mapred.Task: Task 'attempt_local1633189946_0001_m_000000_0' done.
2026-04-21 13:59:07,840 INFO mapred.Task: Final counters for attempt_local1633189946_0001_m_000000_0: Counters: 23
File System Counters
FILE: Number of bytes read=134
FILE: Number of bytes written=630958
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=134
HDFS: Number of bytes written=0
HDFS: Number of read operations=5
HDFS: Number of large read operations=0
HDFS: Number of write operations=1
HDFS: Number of bytes read erasure-coded=0
Map-Reduce Framework
Map input records=6
Map output records=6
Map output bytes=56
Map output materialized bytes=74
Input split bytes=101
Combine input records=0
Spilled Records=6
Failed Shuffles=0
Merged Map outputs=0
GC time elapsed (ms)=112
Total committed heap usage (bytes)=549453824
File Input Format Counters
Bytes Read=134
2026-04-21 13:59:07,840 INFO mapred.LocalJobRunner: Finishing task: attempt_local1633189946_0001_m_000000_0
2026-04-21 13:59:07,841 INFO mapred.LocalJobRunner: Map task executor complete.
2026-04-21 13:59:07,844 INFO mapred.LocalJobRunner: Waiting for reduce tasks
2026-04-21 13:59:07,844 INFO mapred.LocalJobRunner: Starting task: attempt_local1633189946_0001_r_000000_0
2026-04-21 13:59:07,853 INFO output.PathOutputCommitterFactory: No output committer factory defined, defaulting to FileOutputCommitterFactory
2026-04-21 13:59:07,853 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2026-04-21 13:59:07,853 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false
2026-04-21 13:59:07,854 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
2026-04-21 13:59:07,857 INFO reduce.ShuffleTask: Using ShuffleConsumerPlugin: org.apache.hadoop.mapreduce.task.reduce.Shuffle@71479053
2026-04-21 13:59:07,859 WARN Impl.MetricsSystemImpl: JobTracker metrics system already initialized!
2026-04-21 13:59:07,879 INFO reduce.MergeManagerImpl: MergeManager: memoryLimit=2589196288, maxSingleShuffleLimit=647299072, mergeThreshold=1708809632, ioSortFactor=10, memToMemMergeOutputsThreshold=10
2026-04-21 13:59:07,882 INFO reduce.EventFetcher: attempt_local1633189946_0001_r_000000_0 Thread started: EventFetcher for fetching Map Completion Events
```



```
Activities Terminal Apr 21 13:59 lab5@lab5-Latitude-7490: ~/Desktop/Hadoop_12

Total committed heap usage (bytes)=549453824
File Input Format Counters
  Bytes Read=134
2026-04-21 13:59:07,840 INFO mapred.LocalJobRunner: Finishing task: attempt_local1633189946_0001_r_000000_0
2026-04-21 13:59:07,841 INFO mapred.LocalJobRunner: map task executor complete.
2026-04-21 13:59:07,844 INFO mapred.LocalJobRunner: Waiting for reduce tasks
2026-04-21 13:59:07,844 INFO mapred.LocalJobRunner: Starting task: attempt_local1633189946_0001_r_000000_0
2026-04-21 13:59:07,853 INFO output.FileOutputCommitterFactory: No output committer factory defined, defaulting to FileOutputCommitterFactory
2026-04-21 13:59:07,853 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2026-04-21 13:59:07,853 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup_temporary folders under output directory:false, ignore cleanup failures: false
2026-04-21 13:59:07,854 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
2026-04-21 13:59:07,857 INFO mapred.ReduceTask: Using ShuffleConsumerPlugin: org.apache.hadoop.mapreduce.task.reduce.Shuffle@71479653
2026-04-21 13:59:07,859 WARN Inpl.MetricsSystemImpl: JobTracker metrics system already initialized!
2026-04-21 13:59:07,879 INFO reduce.MergeManagerImpl: MergeManager: memoryLimit=2589106280, maxShuffleShuffleLimit=647299072, mergeThreshold=1768869632, toSortFactor=10, memToMemMergeOutputsThreshold=10
2026-04-21 13:59:07,882 INFO reduce.EventFetcher: attempt_local1633189946_0001_r_000000_0 Thread started: EventFetcher for fetching Map completion Events
2026-04-21 13:59:07,917 INFO reduce.LocalFetcher: localFetcher#1 about to shuffle output of map attempt_local1633189946_0001_m_000000_0 decomp: 70 len: 74 to MEMORY
2026-04-21 13:59:07,920 INFO reduce.InMemoryMapOutput: Read 70 bytes from map-output for attempt_local1633189946_0001_m_000000_0
2026-04-21 13:59:07,922 INFO reduce.MergeManagerImpl: closeInMemoryFile -> map-output of size: 70, inMemoryMapOutputs.size() -> 1, commitMemory -> 0, usedMemory -> 70
2026-04-21 13:59:07,924 INFO reduce.EventFetcher: EventFetcher is interrupted.. Returning
2026-04-21 13:59:07,924 INFO mapred.LocalJobRunner: 1 / 1 copied.
2026-04-21 13:59:07,925 INFO reduce.MergeManagerImpl: finalMerge called with 1 in-memory map-outputs and 0 on-disk map-outputs
2026-04-21 13:59:07,931 INFO mapred.Merger: Merging 1 sorted segments
2026-04-21 13:59:07,932 INFO mapred.Merger: Down to the last merge-pass, with 1 segments left of total size: 62 bytes
2026-04-21 13:59:07,933 INFO reduce.MergeManagerImpl: Merged 1 segments, 70 bytes to disk to satisfy reduce memory limit
2026-04-21 13:59:07,934 INFO reduce.MergeManagerImpl: Merging 1 files, 74 bytes from disk
2026-04-21 13:59:07,935 INFO reduce.MergeManagerImpl: Merging 0 segments, 0 bytes from memory into reduce
2026-04-21 13:59:07,935 INFO mapred.Merger: Merging 1 sorted segments
2026-04-21 13:59:07,936 INFO mapred.Merger: Down to the last merge-pass, with 1 segments left of total size: 62 bytes
2026-04-21 13:59:07,936 INFO mapred.LocalJobRunner: 1 / 1 copied.
2026-04-21 13:59:07,974 INFO Configuration.deprecation: mapred.skip.on is deprecated. Instead, use mapreduce.job.skiprecords
2026-04-21 13:59:08,040 INFO mapred.Task: Task:attempt_local1633189946_0001_r_000000_0 is done. And is in the process of committing
2026-04-21 13:59:08,043 INFO mapred.LocalJobRunner: 1 / 1 copied.
2026-04-21 13:59:08,044 INFO mapred.Task: Task:attempt_local1633189946_0001_r_000000_0 is allowed to commit now
2026-04-21 13:59:08,049 INFO output.FileOutputCommitter: Saved output of task 'attempt_local1633189946_0001_r_000000_0' to hdfs://localhost:9000/output
2026-04-21 13:59:08,060 INFO mapred.LocalJobRunner: reduce > reduce
2026-04-21 13:59:08,060 INFO mapred.Task: Task 'attempt_local1633189946_0001_r_000000_0' done.
2026-04-21 13:59:08,061 INFO mapred.Task: Final Counters for attempt_local1633189946_0001_r_000000_0: Counters: 30
File System Counters
  FILE: Number of bytes read=1442
  FILE: Number of bytes written=639032
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=134
  HDFS: Number of bytes written=22
  HDFS: Number of read operations=10
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=3
  HDFS: Number of bytes read erasure-coded=0
Map-Reduce Framework
  Combine input records=0
  Combine output records=0
  Reduce input groups=3
  Reduce shuffle bytes=74
  Reduce input records=6
  Reduce output records=3
  Spilled Records=6
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=0
Total committed heap usage (bytes)=549453824
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Output Format Counters
  Bytes Written=22
2026-04-21 13:59:08,061 INFO mapred.LocalJobRunner: Finishing task: attempt_local1633189946_0001_r_000000_0
2026-04-21 13:59:08,061 INFO mapred.LocalJobRunner: reduce task executor complete.
2026-04-21 13:59:08,418 INFO mapreduce.Job: Job job_local1633189946_0001 running in uber mode : false
2026-04-21 13:59:08,420 INFO mapreduce.Job: map 100% reduce 100%
```

```
Activities Terminal Apr 21 13:59 lab5@lab5-Latitude-7490: ~/Desktop/Hadoop_12

2026-04-21 13:59:07,934 INFO reduce.MergeManagerImpl: Merging 1 files, 74 bytes from disk
2026-04-21 13:59:07,935 INFO reduce.MergeManagerImpl: Merging 0 segments, 0 bytes from memory into reduce
2026-04-21 13:59:07,935 INFO mapred.Merger: Merging 1 sorted segments
2026-04-21 13:59:07,936 INFO mapred.Merger: Down to the last merge-pass, with 1 segments left of total size: 62 bytes
2026-04-21 13:59:07,936 INFO mapred.LocalJobRunner: 1 / 1 copied.
2026-04-21 13:59:07,974 INFO Configuration.deprecation: mapred.skip.on is deprecated. Instead, use mapreduce.job.skiprecords
2026-04-21 13:59:08,040 INFO mapred.Task: Task:attempt_local1633189946_0001_r_000000_0 is done. And is in the process of committing
2026-04-21 13:59:08,043 INFO mapred.LocalJobRunner: 1 / 1 copied.
2026-04-21 13:59:08,044 INFO mapred.Task: Task:attempt_local1633189946_0001_r_000000_0 is allowed to commit now
2026-04-21 13:59:08,049 INFO output.FileOutputCommitter: Saved output of task 'attempt_local1633189946_0001_r_000000_0' to hdfs://localhost:9000/output
2026-04-21 13:59:08,060 INFO mapred.LocalJobRunner: reduce > reduce
2026-04-21 13:59:08,060 INFO mapred.Task: Task 'attempt_local1633189946_0001_r_000000_0' done.
2026-04-21 13:59:08,061 INFO mapred.Task: Final Counters for attempt_local1633189946_0001_r_000000_0: Counters: 30
File System Counters
  FILE: Number of bytes read=1442
  FILE: Number of bytes written=639032
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=134
  HDFS: Number of bytes written=22
  HDFS: Number of read operations=10
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=3
  HDFS: Number of bytes read erasure-coded=0
Map-Reduce Framework
  Combine input records=0
  Combine output records=0
  Reduce input groups=3
  Reduce shuffle bytes=74
  Reduce input records=6
  Reduce output records=3
  Spilled Records=6
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=0
Total committed heap usage (bytes)=549453824
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Output Format Counters
  Bytes Written=22
2026-04-21 13:59:08,061 INFO mapred.LocalJobRunner: Finishing task: attempt_local1633189946_0001_r_000000_0
2026-04-21 13:59:08,061 INFO mapred.LocalJobRunner: reduce task executor complete.
2026-04-21 13:59:08,418 INFO mapreduce.Job: Job job_local1633189946_0001 running in uber mode : false
2026-04-21 13:59:08,420 INFO mapreduce.Job: map 100% reduce 100%
```

```
Activities Terminal Apr 21 13:59 lab5@lab5-Latitude-7490: ~/Desktop/Hadoop_12

GC time elapsed (ms)=0
Total committed heap usage (bytes)=549453824
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Output Format Counters
  Bytes Written=22
2026-04-21 13:59:08,061 INFO mapred.LocalJobRunner: Finishing task: attempt local1633189946_0001_r_000000_0
2026-04-21 13:59:08,061 INFO mapred.LocalJobRunner: reduce task executor complete.
2026-04-21 13:59:08,418 INFO mapreduce.Job: Job job_local1633189946_0001 running in uber mode : false
2026-04-21 13:59:08,420 INFO mapreduce.Job: map 100% reduce 100%
2026-04-21 13:59:08,423 INFO mapreduce.Job: Job job_local1633189946_0001 completed successfully
2026-04-21 13:59:08,442 INFO mapreduce.Job: Counters: 36
File System Counters
  FILE: Number of bytes read=6704
  FILE: Number of bytes written=1277990
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=208
  HDFS: Number of bytes written=22
  HDFS: Number of read operations=15
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=4
  HDFS: Number of bytes read erasure-coded=0
Map-Reduce Framework
  Map input records=6
  Map output records=6
  Map output bytes=56
  Map output materialized bytes=74
  Input split bytes=101
  Combine input records=0
  Combine output records=0
  Reduce input groups=3
  Reduce shuffle bytes=74
  Reduce input records=6
  Reduce output records=3
  Spilled Records=12
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=112
  Total committed heap usage (bytes)=1098907648
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
```

```
Activities Terminal Apr 21 14:00 lab5@lab5-Latitude-7490: ~/Desktop/Hadoop_12

2026-04-21 13:59:08,061 INFO mapred.LocalJobRunner: reduce task executor complete.
2026-04-21 13:59:08,418 INFO mapreduce.Job: Job job_local1633189946_0001 running in uber mode : false
2026-04-21 13:59:08,420 INFO mapreduce.Job: map 100% reduce 100%
2026-04-21 13:59:08,423 INFO mapreduce.Job: Job job_local1633189946_0001 completed successfully
2026-04-21 13:59:08,442 INFO mapreduce.Job: Counters: 36
File System Counters
  FILE: Number of bytes read=6704
  FILE: Number of bytes written=1277990
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=208
  HDFS: Number of bytes written=22
  HDFS: Number of read operations=15
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=4
  HDFS: Number of bytes read erasure-coded=0
Map-Reduce Framework
  Map input records=6
  Map output records=6
  Map output bytes=56
  Map output materialized bytes=74
  Input split bytes=101
  Combine input records=0
  Combine output records=0
  Reduce input groups=3
  Reduce shuffle bytes=74
  Reduce input records=6
  Reduce output records=3
  Spilled Records=12
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=112
  Total committed heap usage (bytes)=1098907648
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=134
File Output Format Counters
  Bytes Written=22
lab5@lab5-Latitude-7490:~/Desktop/Hadoop_12$ hdfs dfs -cat /output/part-r-00000
ERROR 2
INFO 3
WARN 1
lab5@lab5-Latitude-7490:~/Desktop/Hadoop_12$
```

Practical No. 13

input.txt:

hello spark

hello scala

hello hadoop

apache spark hello

OUTPUT:

```
Activities Terminal Apr 22 14:23 lab5@lab5-Latitude-7490: ~/Desktop/Hadoop13

lab5@lab5-Latitude-7490:~/Desktop/Hadoop13$ spark-shell
26/04/22 14:15:09 WARN Utils: Your hostname, lab5-Latitude-7490 resolves to a loopback address: 127.0.1.1; using 10.95.95.65 instead (on interface wlp2s0)
26/04/22 14:15:09 WARN Utils: Set SPARK_LOCAL_IP if you need to bind to another address
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
26/04/22 14:15:14 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Spark context Web UI available at http://10.95.95.65:4040
Spark context available as 'sc' (master = local[*], app id = local-1776847515803).
Spark session available as 'spark'.
Welcome to

  ____      __
 / ___ |    /  \
/ /___|    /    \
\___  |   /  /\
     |  /____\
     |_/_____\

version 3.5.0

Using Scala version 2.12.18 (OpenJDK 64-Bit Server VM, Java 1.8.0_452)
Type in expressions to have them evaluated.
Type :help for more information.

scala> System.setProperty("user.dir", "/home/lab5/Desktop/Hadoop13")
res0: String = /home/lab5/Desktop/Hadoop13

scala> val file = sc.textFile("input.txt")
file: org.apache.spark.rdd.RDD[String] = input.txt MapPartitionsRDD[1] at textFile at <console>:23

scala> val result = file.flatMap(_.split(" "))
result: org.apache.spark.rdd.RDD[String] = MapPartitionsRDD[2] at FlatMap at <console>:23

scala> .map(word => (word,1))
res1: org.apache.spark.rdd.RDD[(String, Int)] = MapPartitionsRDD[3] at map at <console>:24

scala> .reduceByKey(_ + _)
res2: org.apache.spark.rdd.RDD[(String, Int)] = ShuffledRDD[4] at reduceByKey at <console>:24

scala> res2.collect.foreach(println)
(scala,1)
(apache,1)
(hello,4)
(spark,2)
(hadoop,1)

scala> 
```